

# Entrepôts de données

Federico Ulliana

# **QUERY OPTIMIZATION FOR ON-LINE ANALYTICS (OLAP)**

# DW Need Optimization !

- At this stage of the DW design we should have a good model to support information extraction
  - at least one problem is addressed.
- Next step is to deal with data that may be so voluminous (OR queries so resource-demanding) that response times take minutes or hours

# OLAP

## ON-LINE ANALYTICAL QUERY PROCESSING

Although queries may touch millions of rows (= ANALYTICAL),  
the system should answer within seconds (= ON-LINE).

- We will see some interesting optimizations
  - Materialized Views
  - Bitmap Indexes
  - Join Indexes

# Snapshot Book Sales

**Book\_Dim**

Book	Genre
Winnie	Children
C	Inform.

**Fact\_Book\_Sales**

Book	Shop	Date	TotPrice
Winnie	854	11/11	20
C	876	12/11	10
C	854	12/11	5
Winnie	876	13/11	60

- Hypothesis : size fact table  $\sim 10^9$

# Analytics : book selling by genre

```
SELECT Genre, SUM(TotPrice) as Profit  
  
FROM    Fact_Book_Sales F, Book_Dim D  
  
WHERE   F.BookID = D.BookID  
  
GROUP BY D.Genre
```

- Touches all rows in the fact table (can be  $\sim 10^9$ )

# Analytics : book selling by genre

Genre	Profit
Children	1600
C	400
:	

# Materialized Views

- Goal : pre-compute some queries so as to speed up query answering
- Note : the previous query is independent of time

# Materialized Views

**CREATE MATERIALIZED VIEW**

**Salesview (BookID, ShopID, Profit)**

**AS**

SELECT BookID, ShopID, SUM(TotPrice)

FROM Fact\_Book\_Sales

GROUP BY BookID, ShopID

# Materialized Views

Book	Shop	Date	Cnt	TotP
Winnie	854	11/11	1	20
C	876	12/11	2	10
C	854	12/11	1	5
Winnie	876	13/11	3	60
Winnie	954	11/11	1	20
C	976	12/11	2	10
C	994	12/11	1	5
:				

F

# Materialized Views

Book	Shop	Date	Cnt	TotP	Book	Shop	Profit
Winnie	854	11/11	1	20	Winnie	854	600
C	876	12/11	2	10	C	876	200
C	854	12/11	1	5	C	854	200
Winnie	876	13/11	3	60	Winnie	876	1000
Winnie	954	11/11	1	20			
C	976	12/11	2	10			
C	994	12/11	1	5			
:							

F

# Materialized Views

Book	Shop	Date	Cnt	TotP	Book	Shop	Profit
Winnie	854	11/11	1	20	Winnie	854	600
C	876	12/11	2	10	C	876	200
C	854	12/11	1	5	C	854	200
Winnie	876	13/11	3	60	Winnie	876	1000
Winnie	954	11/11	1	20			
C	976	12/11	2	10			
C	994	12/11	1	5			
:							

**F**      **V**

# Materialized Views

Book	Shop	Date	Cnt	TotP	Book	Shop	Profit
Winnie	854	11/11	1	20	Winnie	854	600
C	876	12/11	2	10	C	876	200
C	854	12/11	1	5	C	854	200
Winnie	876	13/11	3	60	Winnie	876	1000
Winnie	954	11/11	1	20			
C	976	12/11	2	10			
C	994	12/11	1	5			
:							

Size Fact Table >> Size View

# Without Materialization Q(**F**)

SELECT      **Genre**, SUM(**TotPrice**)

FROM          **FactSales F**, BookDim D

WHERE        **F**.BookID = D.BookID

GROUP BY D.Genre

# With Materialization

$Q^{\text{ref}}(V)$

SELECT

Genre, SUM(**Profit**)

FROM

**SalesView V, DimBook D**

WHERE

**V.BookID = D.BookID**

GROUP BY D.Genre

# Query Evaluation

Book	Shop	Date	Cnt	TotP
Winnie	854	11/11	1	20
C	876	12/11	2	10
C	854	12/11	1	5
Winnie	876	13/11	3	60
Winnie	954	11/11	1	20
C	976	12/11	2	10
C	994	12/11	1	5
:				

Book	Shop	Profit
Winnie	854	600
C	876	200
C	854	200
Winnie	876	1000

  $Q(F)$

$\gg_{\text{time}}$

  $Q^{\text{ref}}(V)$

# Materialized Views

Pro :

- can reduce query time
- computation done once
- reused by multiple queries

Cons :

- needs storage space
- needs to be maintained
- how to choose the best views ?
  - DB-admin or automatic?
- needs reformulation algorithms

# Query Reformulation Using Views

## Problem

### *Input*

- A (set of) relations  $\mathbf{F}$
- A (set of) views  $\mathbf{V}$
- A (set of) user queries  $\mathbf{Q}$

### *Output*

- $\mathbf{Q}^{\text{ref}}$  such that  $\mathbf{Q}^{\text{ref}}(\mathbf{V}) = \mathbf{Q}(\mathbf{F})$

# View **Relevance** (for Query Reformulation)

- A view has  $\text{sum}(\text{sales})$  by **genre** and by **year** for books introduced after 1991

Genre	Year	SumSales
Children	1991	600
Cooking	1992	800

- is it useful for  $\text{sum}(\text{sales})$  by genre for products introduced after 1992 ?

# View **Relevance** (for Query Reformulation)

- A view has  $\text{sum}(\text{sales})$  by **genre** and by **year** for books introduced after 1991

Genre	Year	SumSales
Children	1991	600
Cooking	1992	800

- is it useful for  $\text{sum}(\text{sales})$  by genre for products introduced before 1991 ?

# View **Relevance** (for Query Reformulation)

- A view has  $\text{sum}(\text{sales})$  by **genre** and by **year** for books introduced after 1991

Genre	Year	SumSales
Children	1991	600
Cooking	1992	800

- is it useful for  $\text{sum}(\text{sales})$  by **region** ?

# View **Relevance** (for Query Reformulation)

- A view has  $\text{sum}(\text{sales})$  by **genre** and by **year** for books introduced after 1991

Genre	Year	SumSales
Children	1991	600
Cooking	1992	800

- is it useful for  $\text{sum}(\text{sales})$  by genre ?

# Query Reformulation

Possible to reformulate a query with a view when :

- The view has selection-conditions that are compatible with those of the queries
- The view has dimensions that are compatible with those of the queries
- In this case, we say that the view is relevant.

# Materialized-View Selection

## Problem

- DW administrator knows typical workload

$$W = \{ Q_1 \dots Q_n \} \quad n=20,50,100$$

of the datawarehouse  $D$ .

Example :

- $Q_1 = \text{SELECT} \dots \text{GROUP BY Product, Date}$
- $Q_2 = \text{SELECT} \dots \text{GROUP BY Product, Shop}$

# Materialized-View Selection

## Problem

- Problem : how to select the **best**-set-of-views  
**V** able to answer **W** (via query reformulation)?
- Unfeasible by hand, needs automatic solution.

# Materialized-View Selection Problem

## Input

- workload  $\mathbf{W}$
- database  $\mathbf{D}$

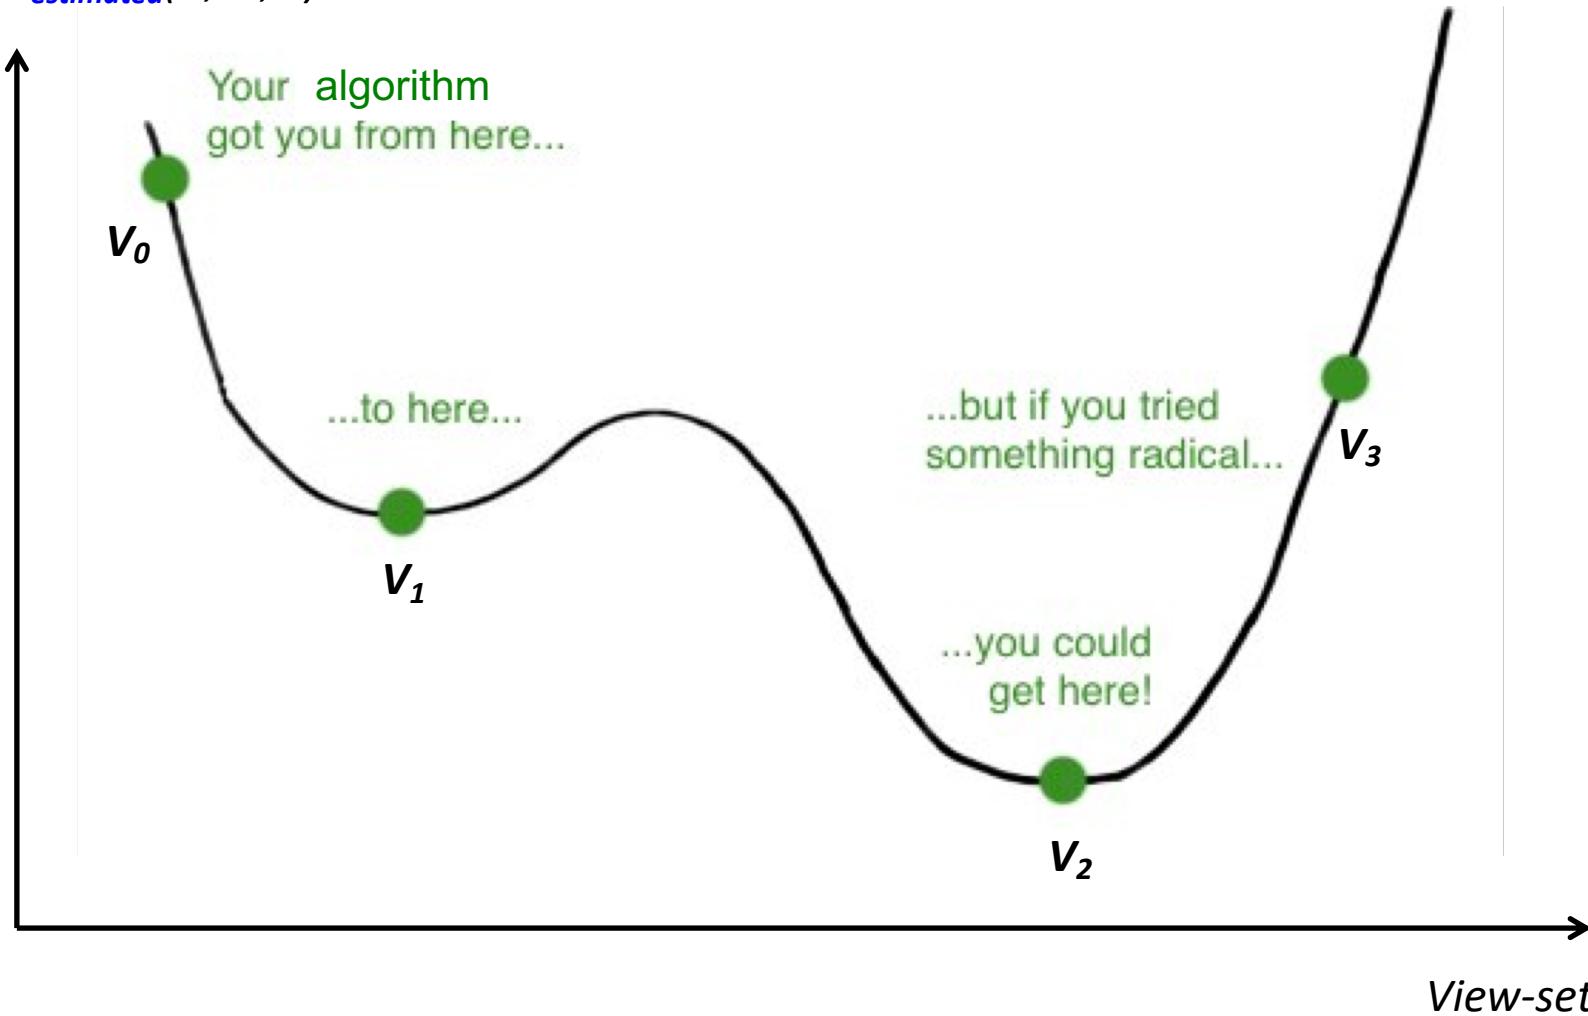
## Output

- A set of views  $\mathbf{V} = \{ V_1 \dots V_m \}$  minimizing  
(a linear combination of)
  1. query reformulation + execution time
  2. view space occupancy
  3. maintenance cost

$$\begin{aligned} & \text{Cost}_{estimated}(\mathbf{V}, \mathbf{W}, \mathbf{D}) \\ &= \alpha_{ref} * \text{reformulation}_{est} \\ &+ \alpha_{occ} * \text{occupancy}_{est} \\ &+ \alpha_{maint} * \text{maintenance}_{est} \end{aligned}$$

# Challenge : cost function is unpredictable

$\text{Cost}_{\text{estimated}}(V, W, D)$



# Materialized View Selection

## Naïve Algorithm

1. For all possible view set  $\mathbf{v}$
2. Compute  $\text{Cost}(\mathbf{v}, \mathbf{w}, \mathbf{D})$
3. Keep  $\mathbf{v}$  with lowest cost

Finds optimal solution, but in an infinite time.

# Materialized View Selection

## Naïve Algorithm

1. For all candidate view set  $\mathbf{v}$
2. Compute  $\text{Cost}(\mathbf{v}, \mathbf{w}, \mathbf{D})$
3. Keep  $\mathbf{v}$  with lowest cost

# Candidate View-Set

## Definition.

A set of views  $V$  is *candidate* for  $W$  and  $D$  if

- (1) it allows to answer all queries in  $W$  by means of reformulation
- (2) every query of  $W$  admit a unique rewriting with  $V$
- (3)  $V$  does not contain “useless” views

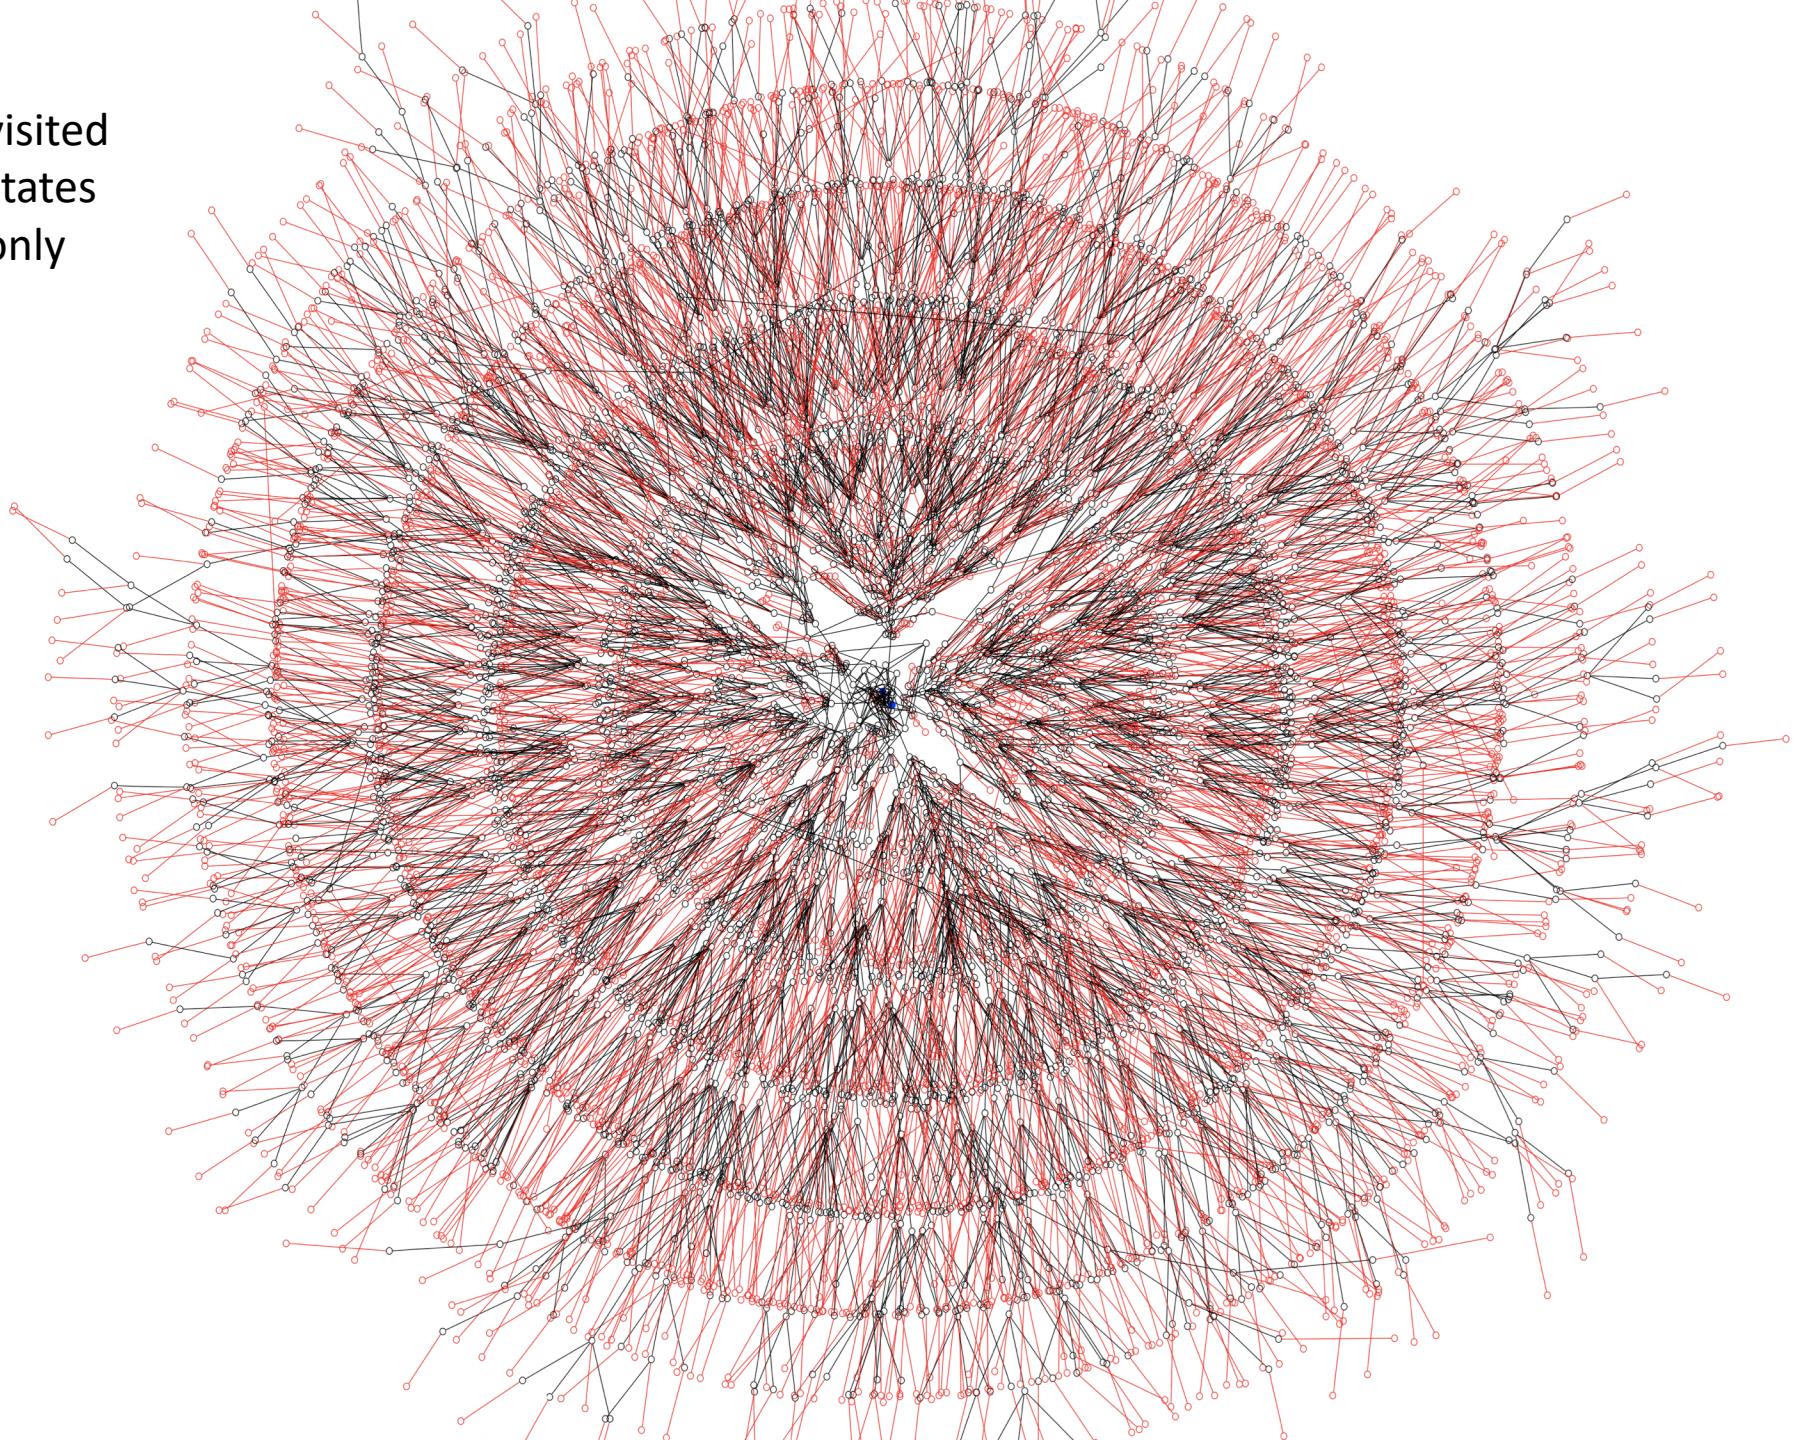
# Materialized View Selection

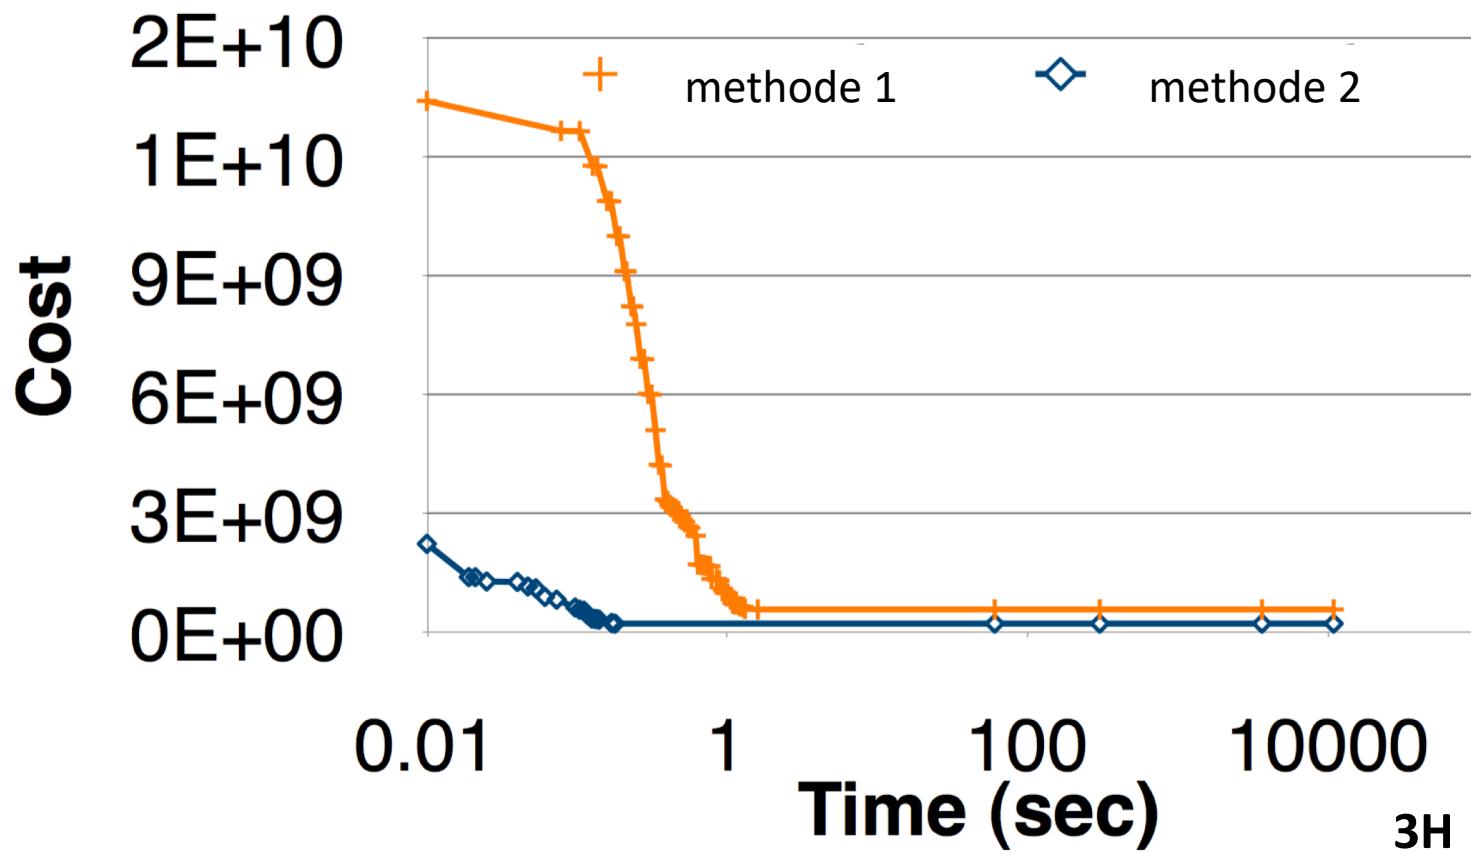
## Naïve Algorithm

1. For all candidate view set  $V$
2. Compute  $\text{Cost}(V, W, D)$
3. Keep  $V$  with lowest cost

We still need to explore a very large space of solutions.

visited  
states  
only



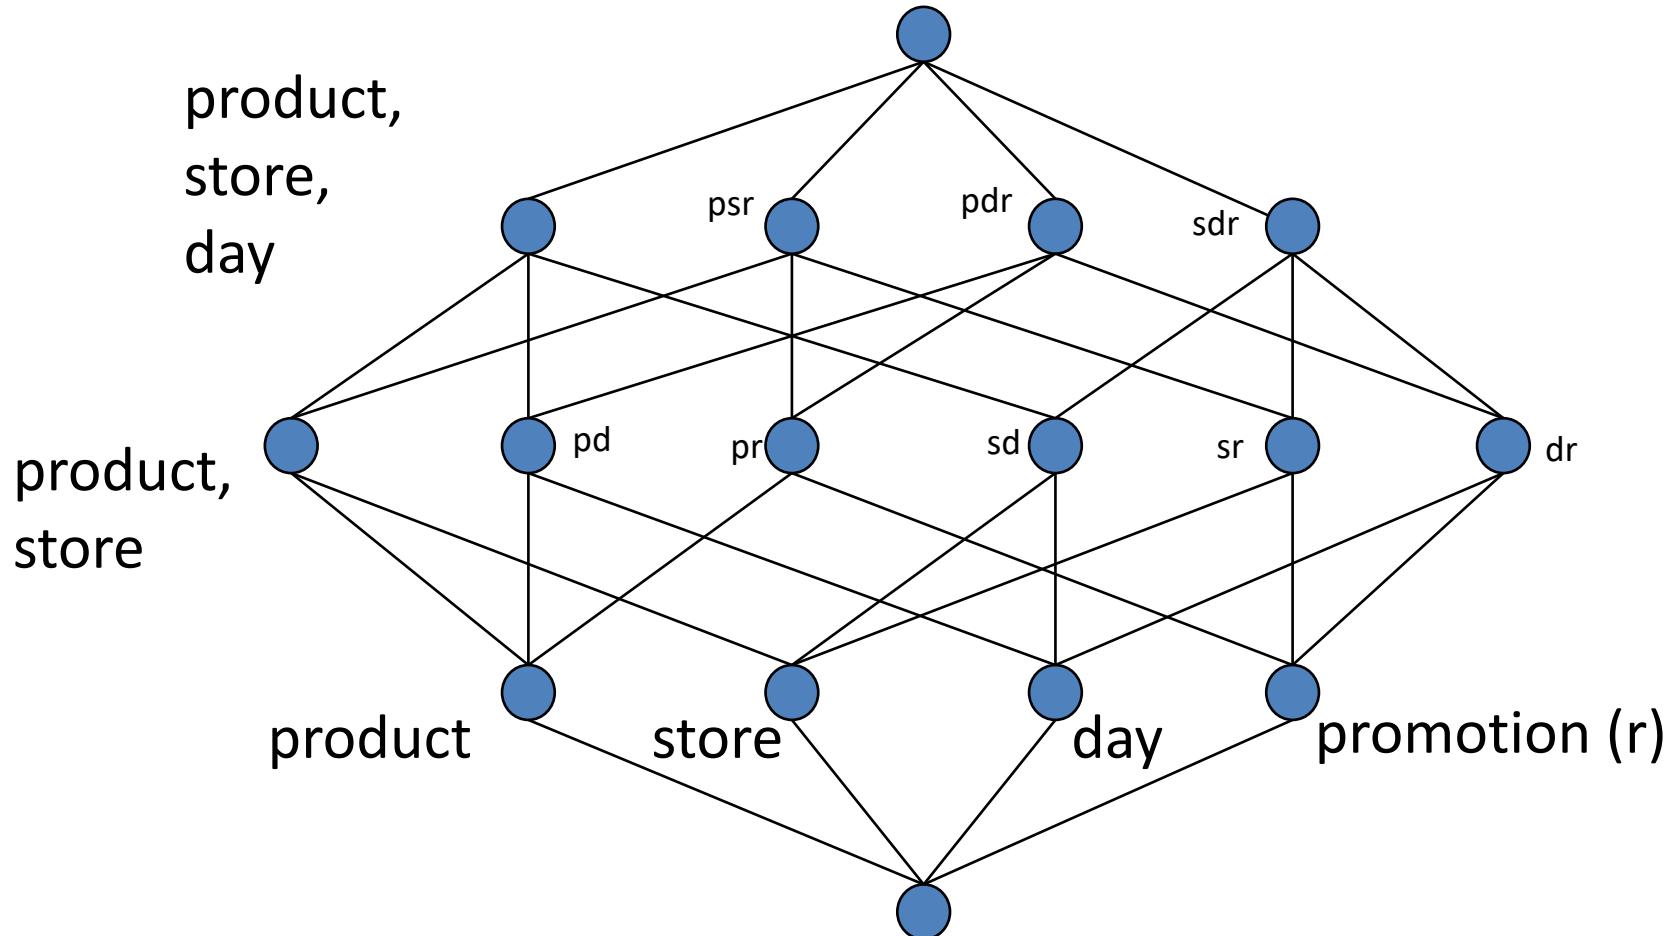


# Solution : split the problem in two

- Consider a simplified view-selection problem where one considers only group-by clauses
- All filtering conditions (like  $\text{year} > 1991$ ) are considered when the former subproblem is solved.

# Aggregation Lattice

product,store,day,promotion (original fact table)

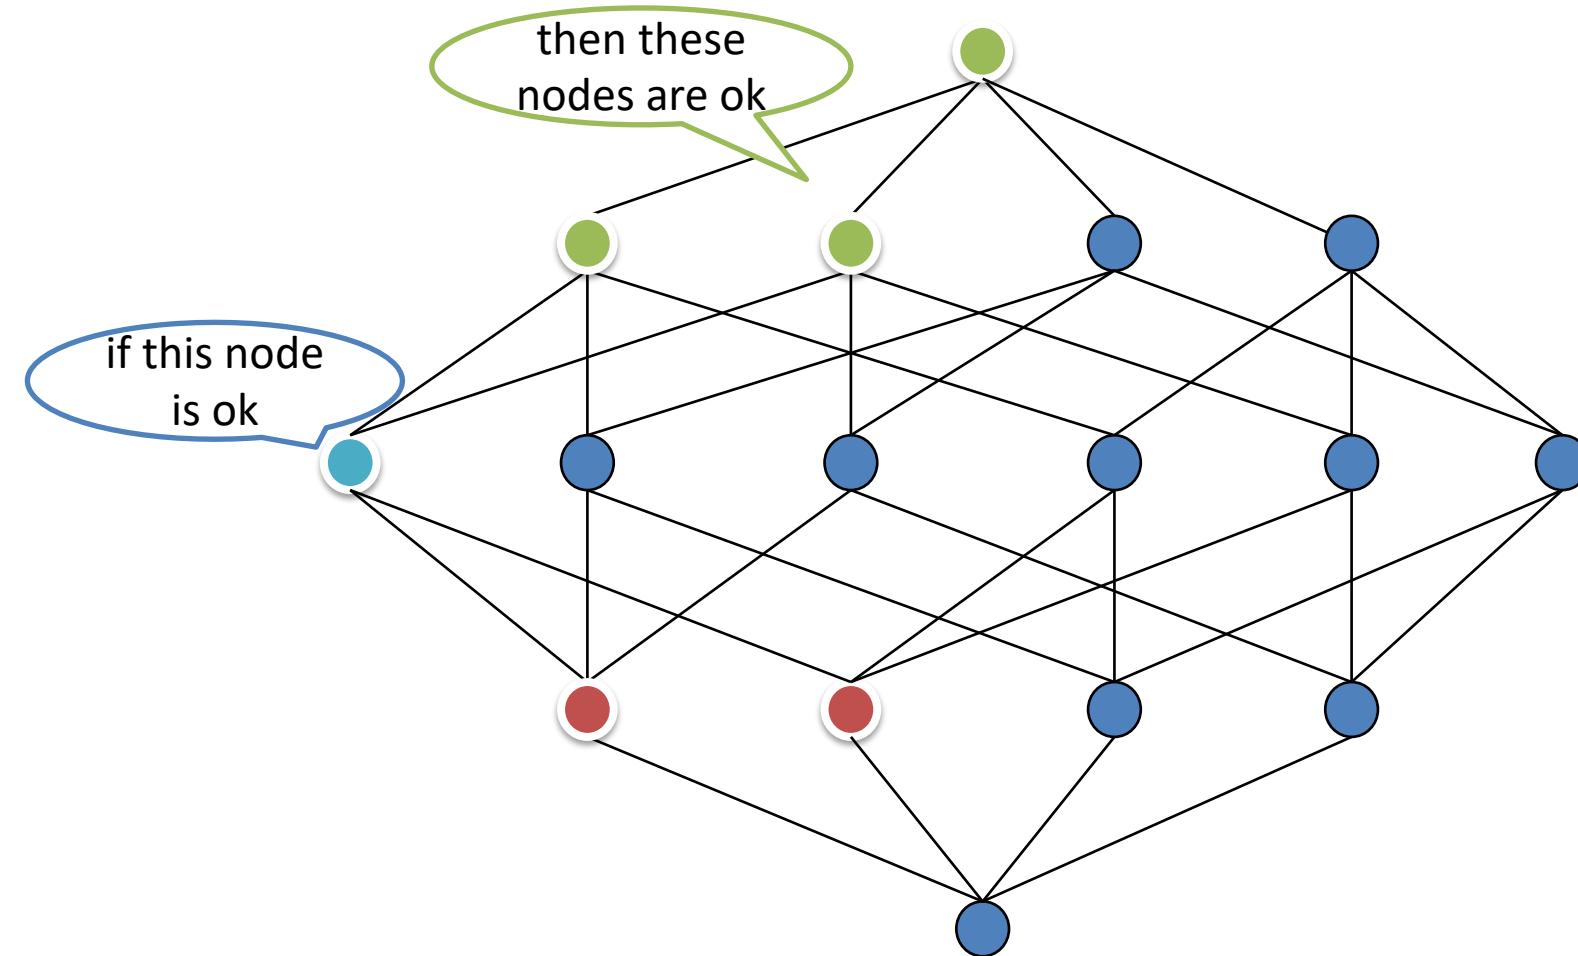


# Aggregation Lattice

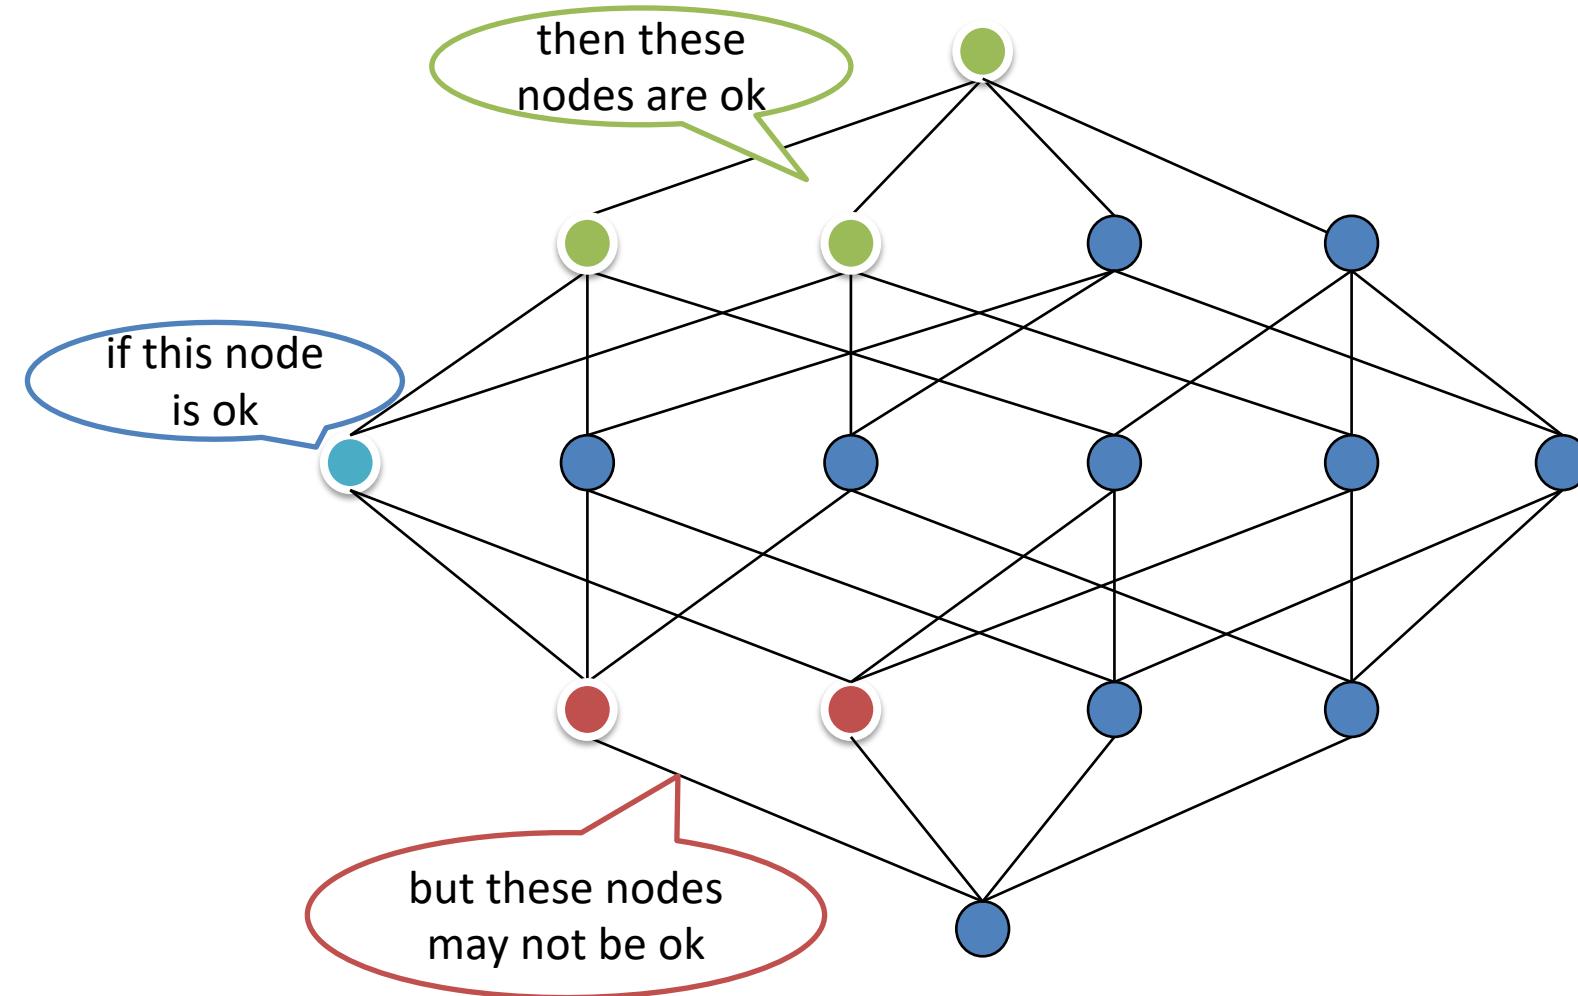
## Property

A group-by query that can be answered by a node **n** of the lattice, can be answered by any ancestor of **n**.

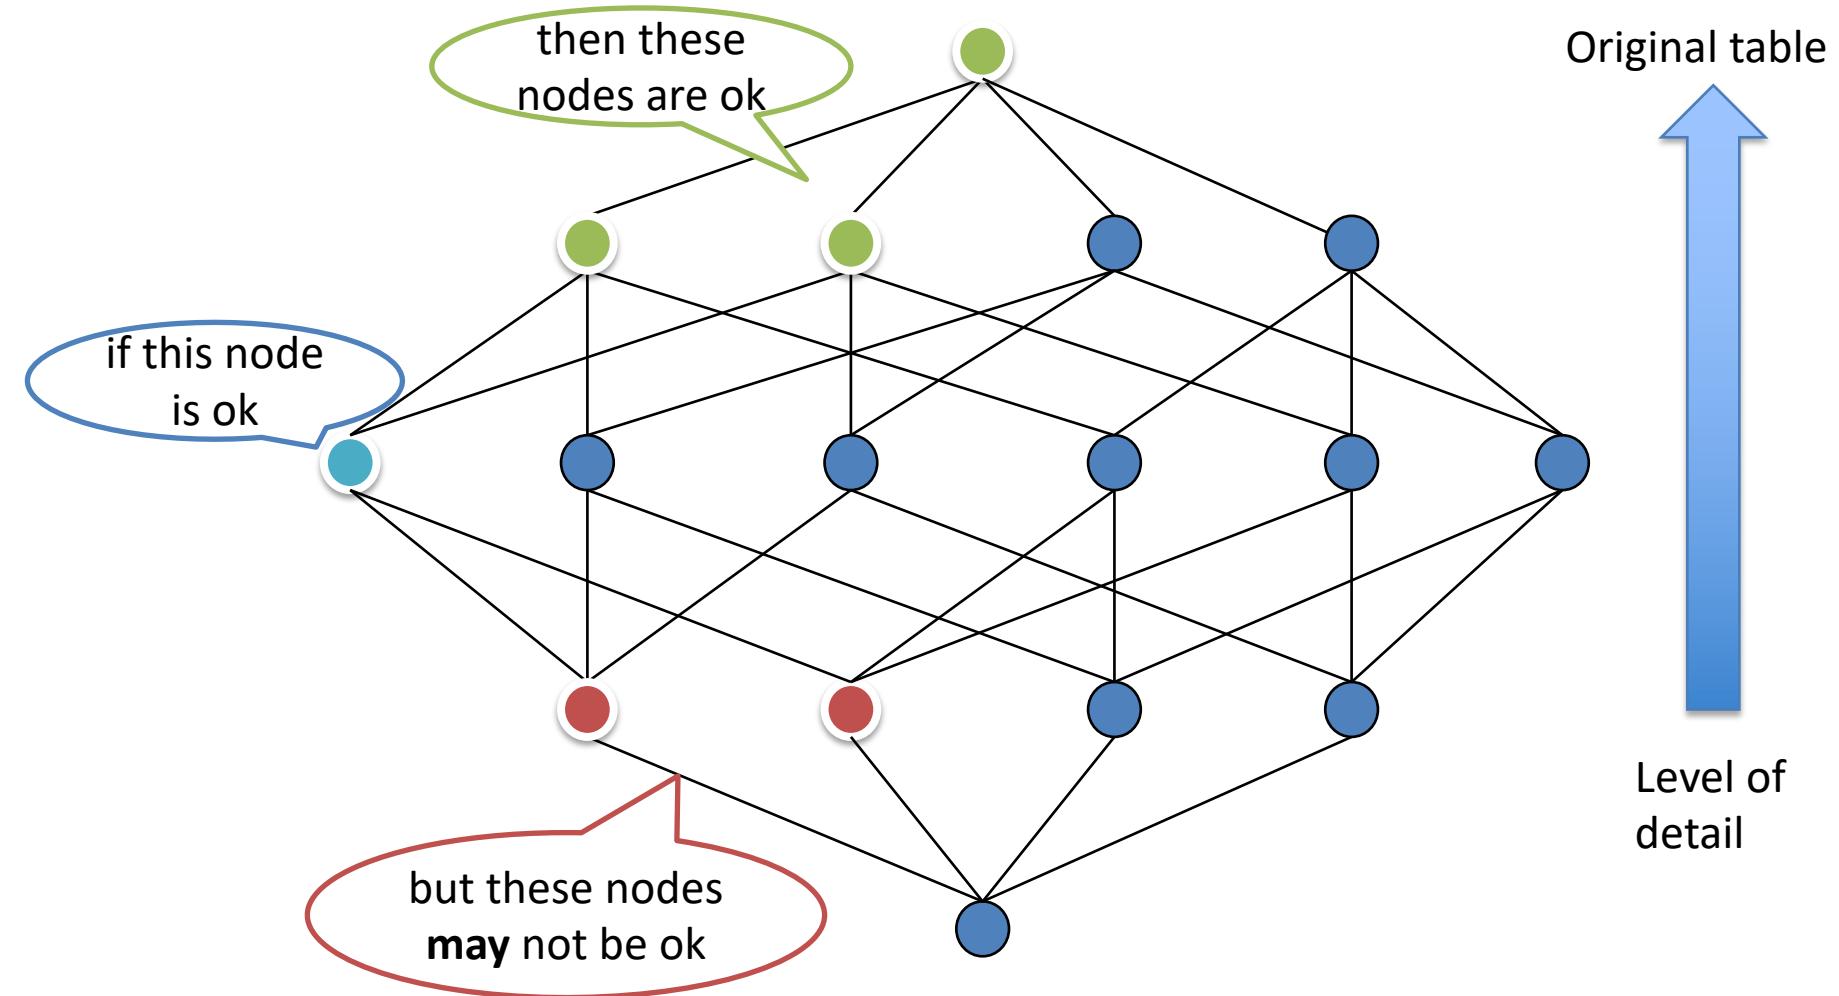
# Aggregation lattice



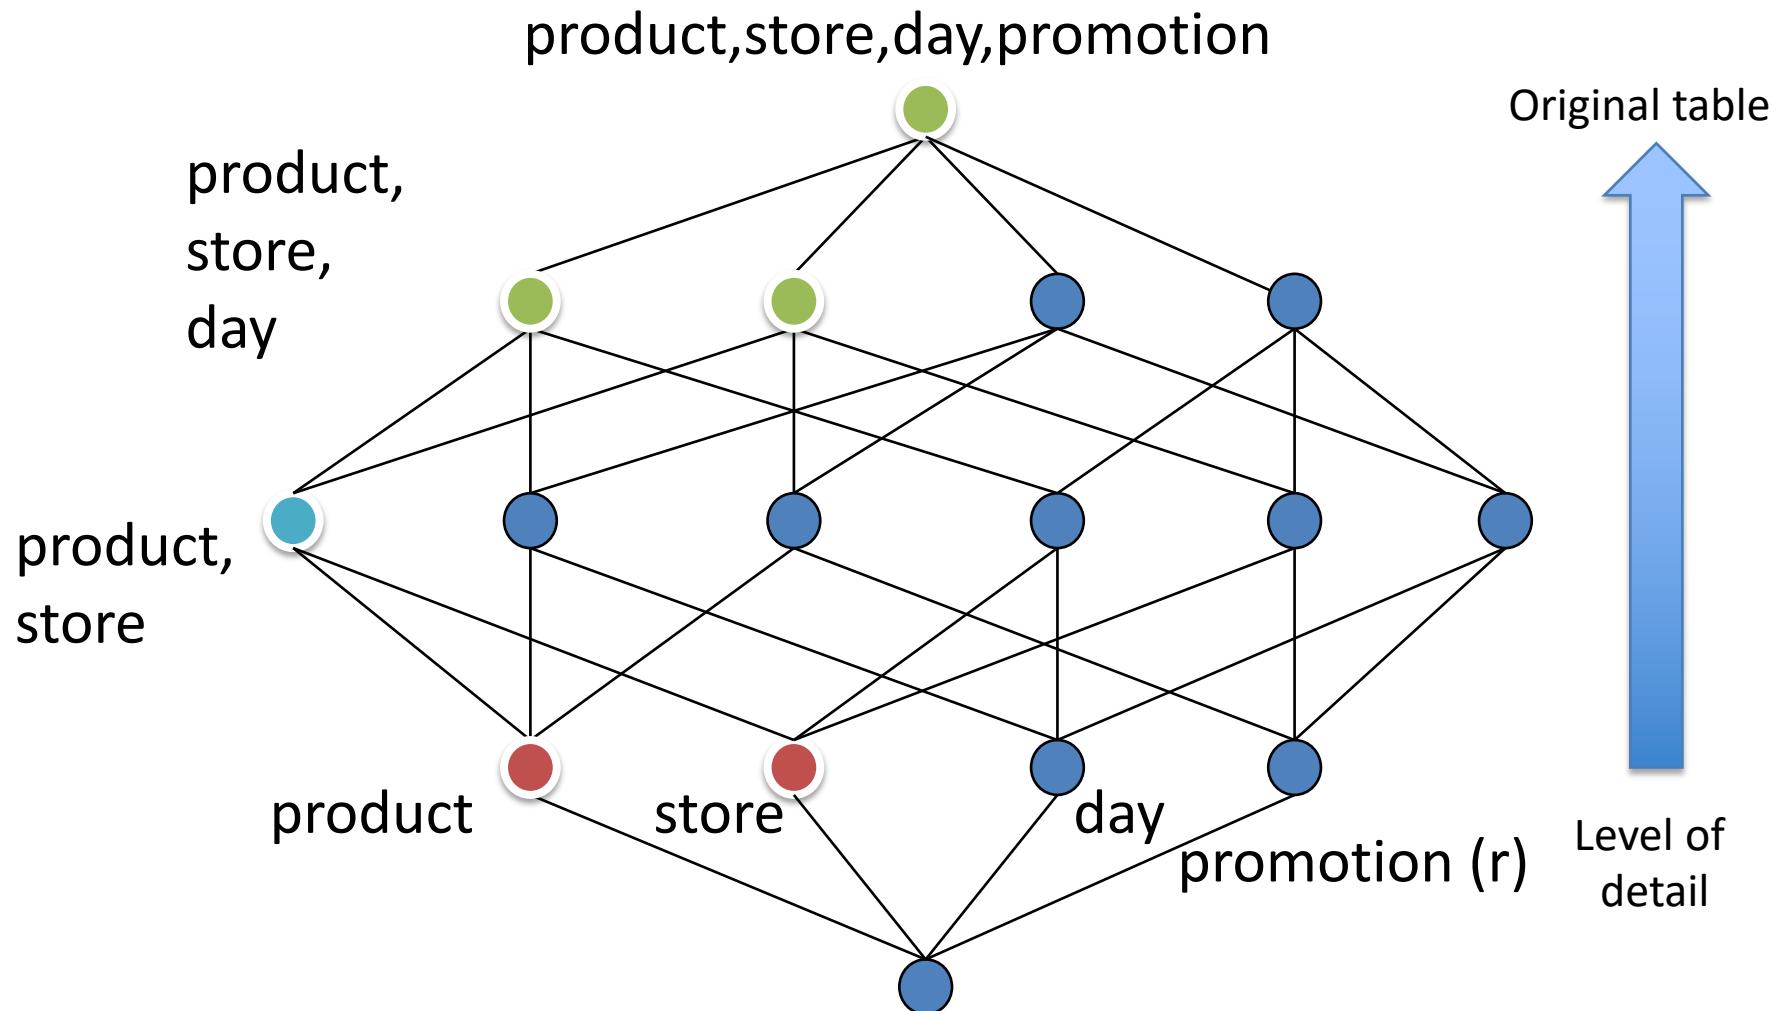
# Aggregation lattice



# Aggregation lattice



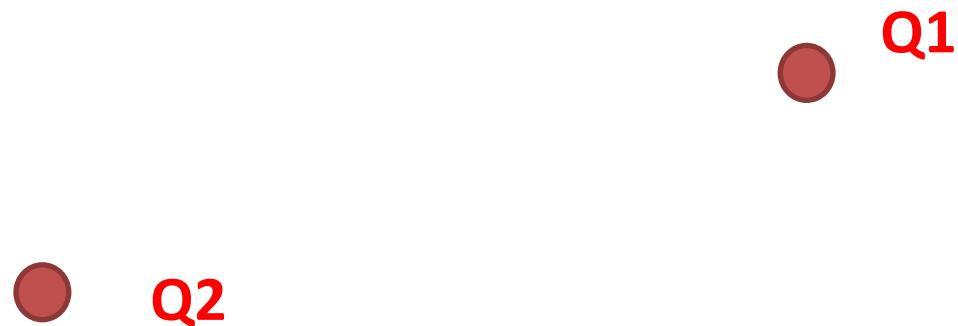
# Aggregation lattice



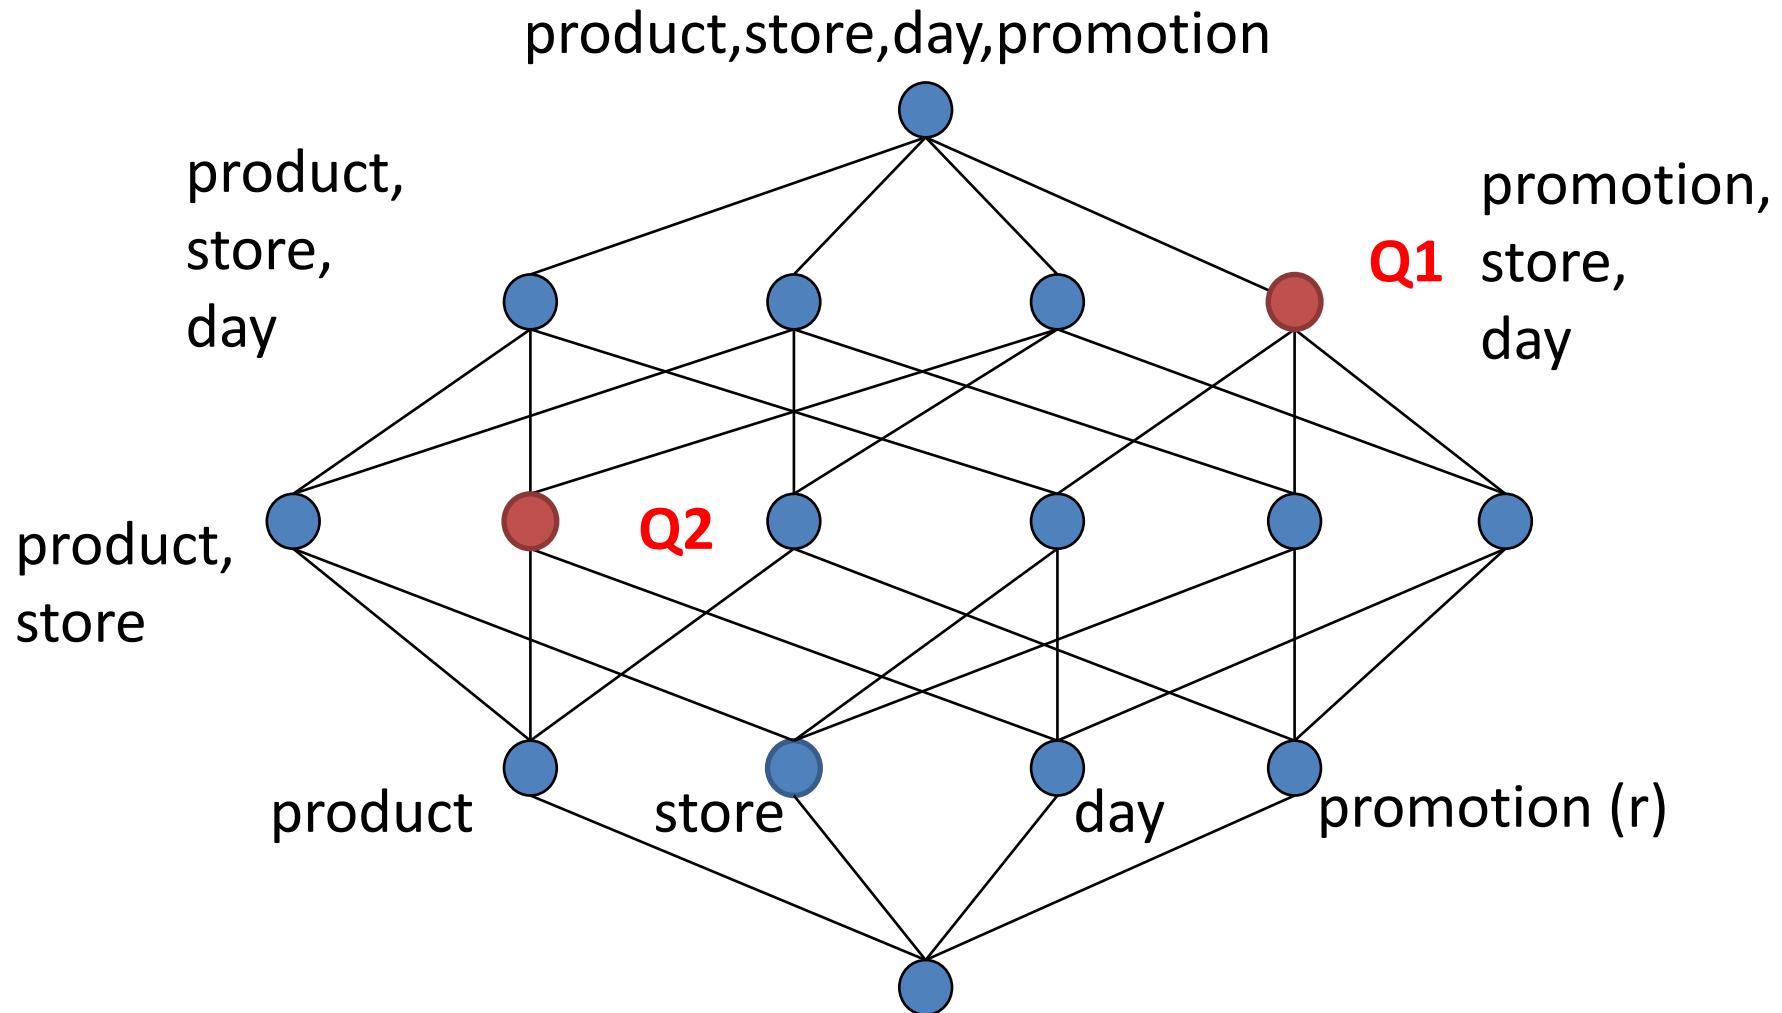
# Materialized-View Selection Algorithm

1. For each  $Q$  in  $W$ , find “lowest” node allowing one to answer  $Q$
2. Prune irrelevant sub-lattice
3. Evaluate cost function on every solution

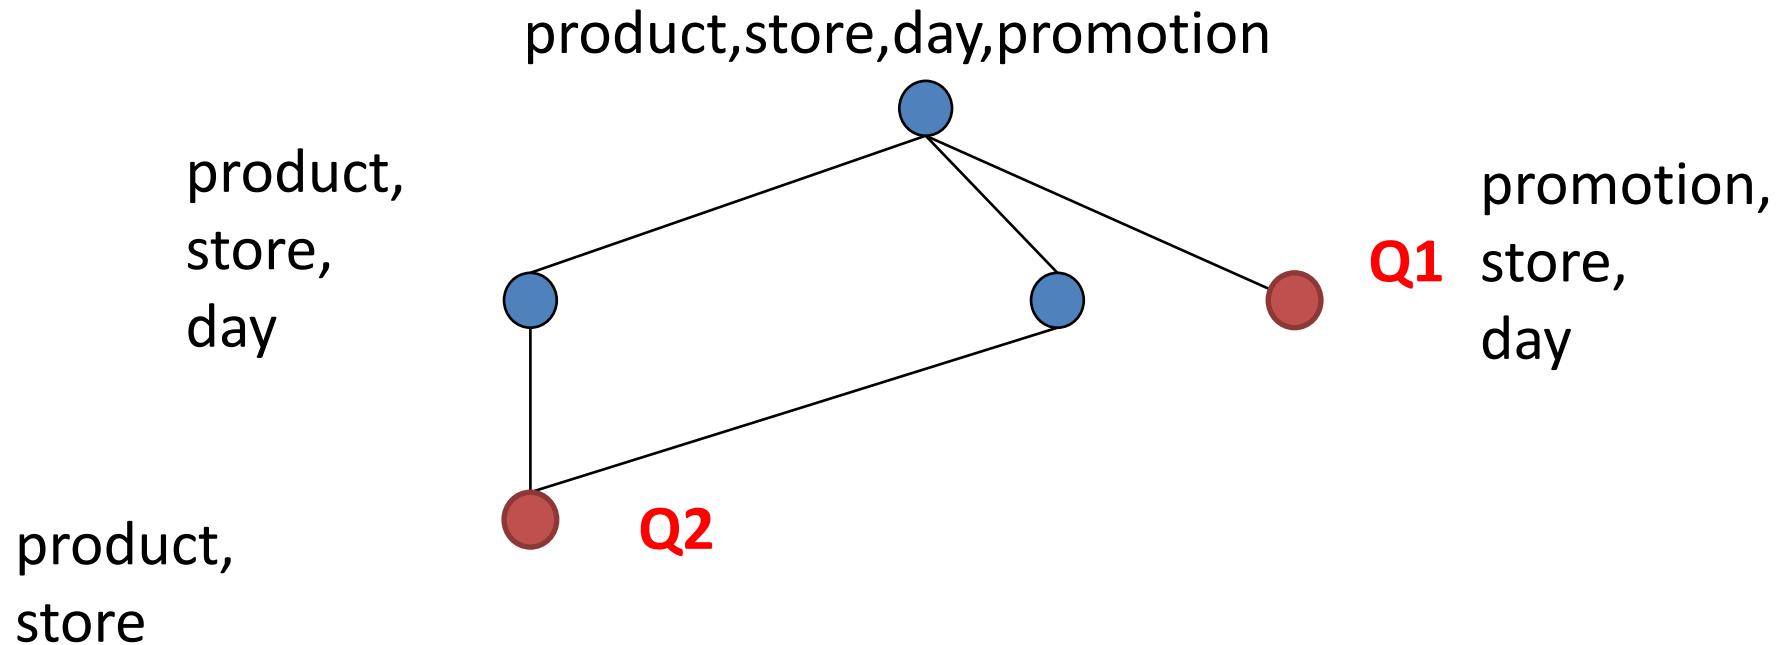
Find lowest node for each Q



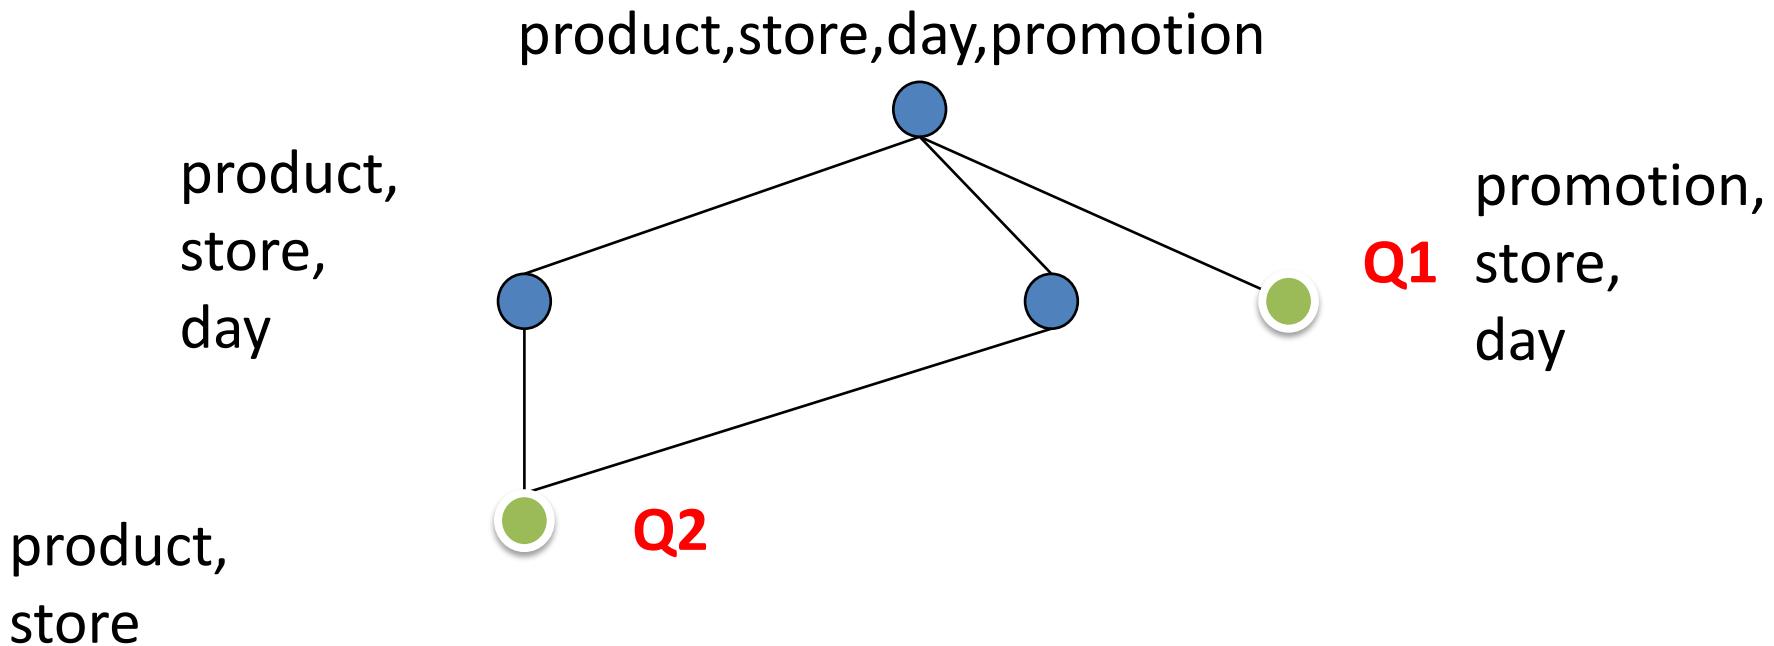
# Find lowest node for each Q



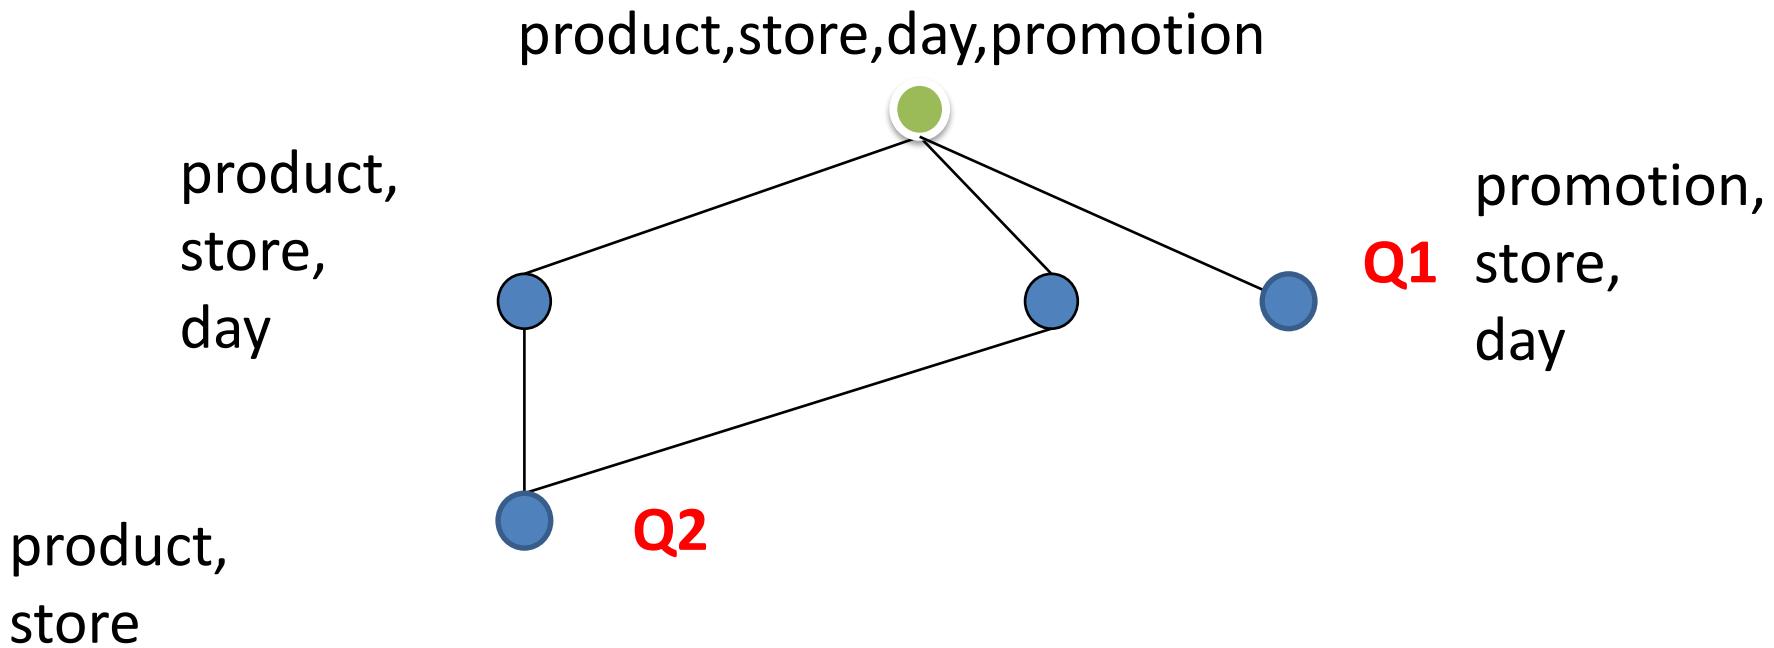
# Prune irrelevant sub-lattice



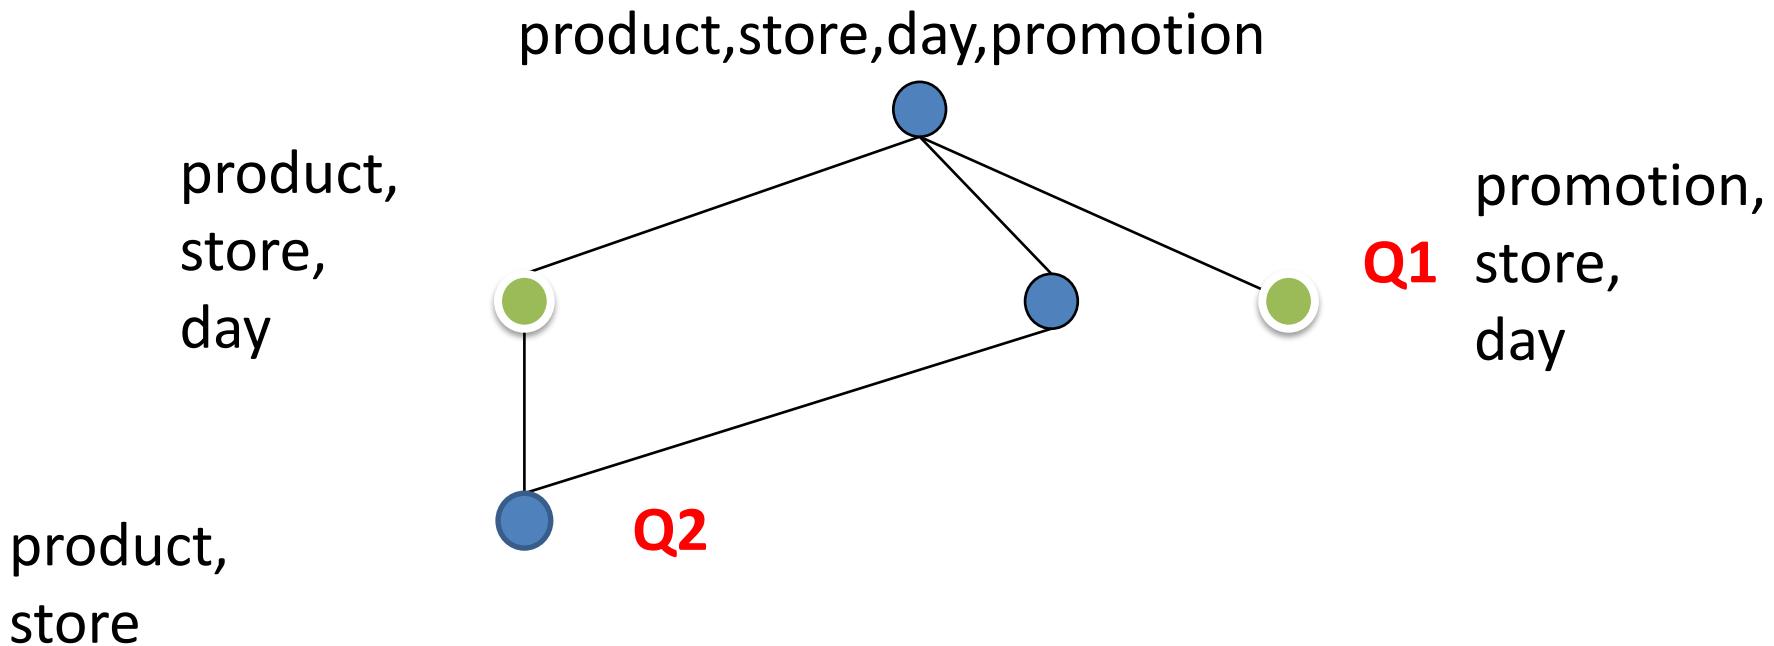
Evaluate cost function on every solution (1)



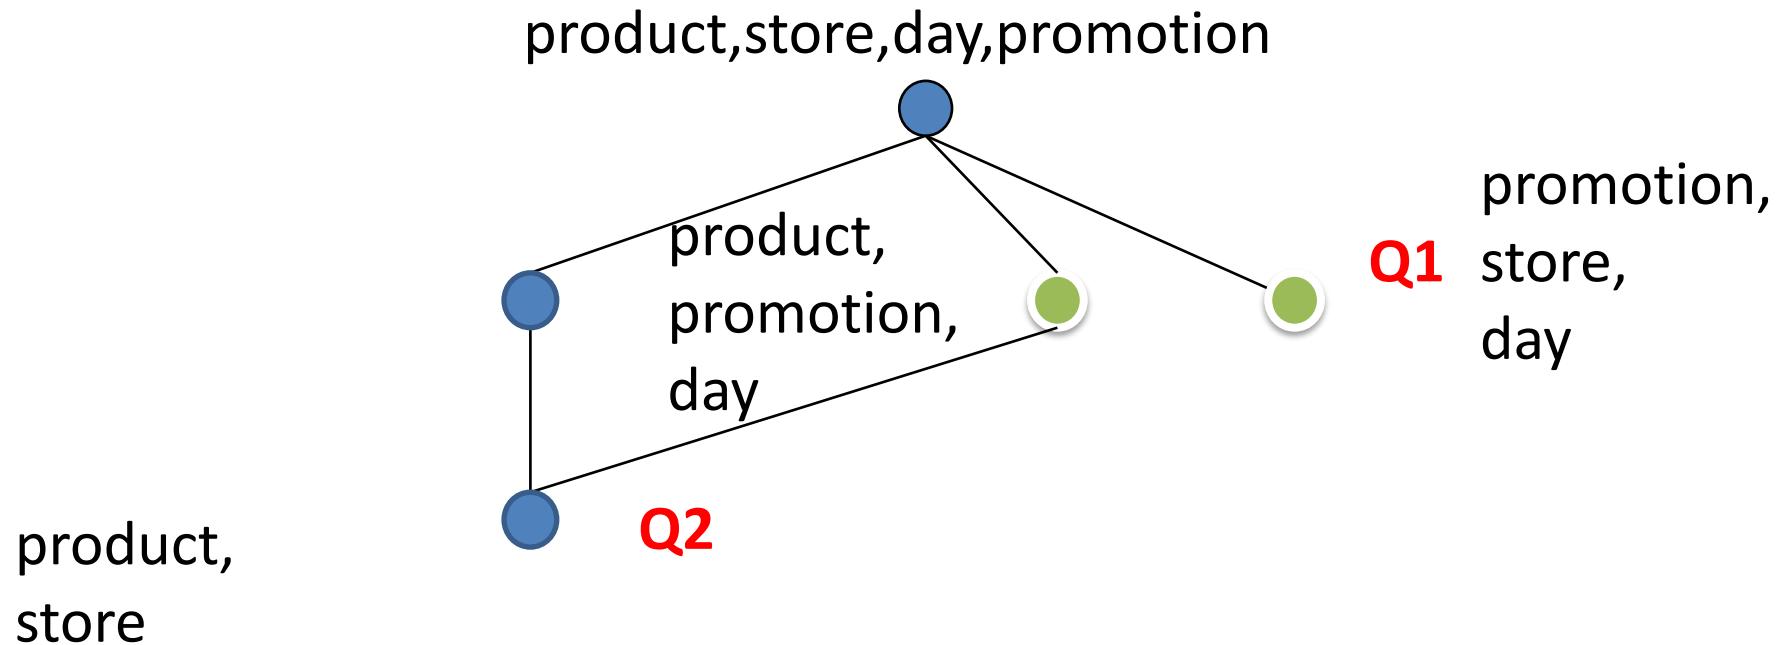
Evaluate cost function on every solution (2)



Evaluate cost function on every solution (3)



Evaluate cost function on every solution (4)

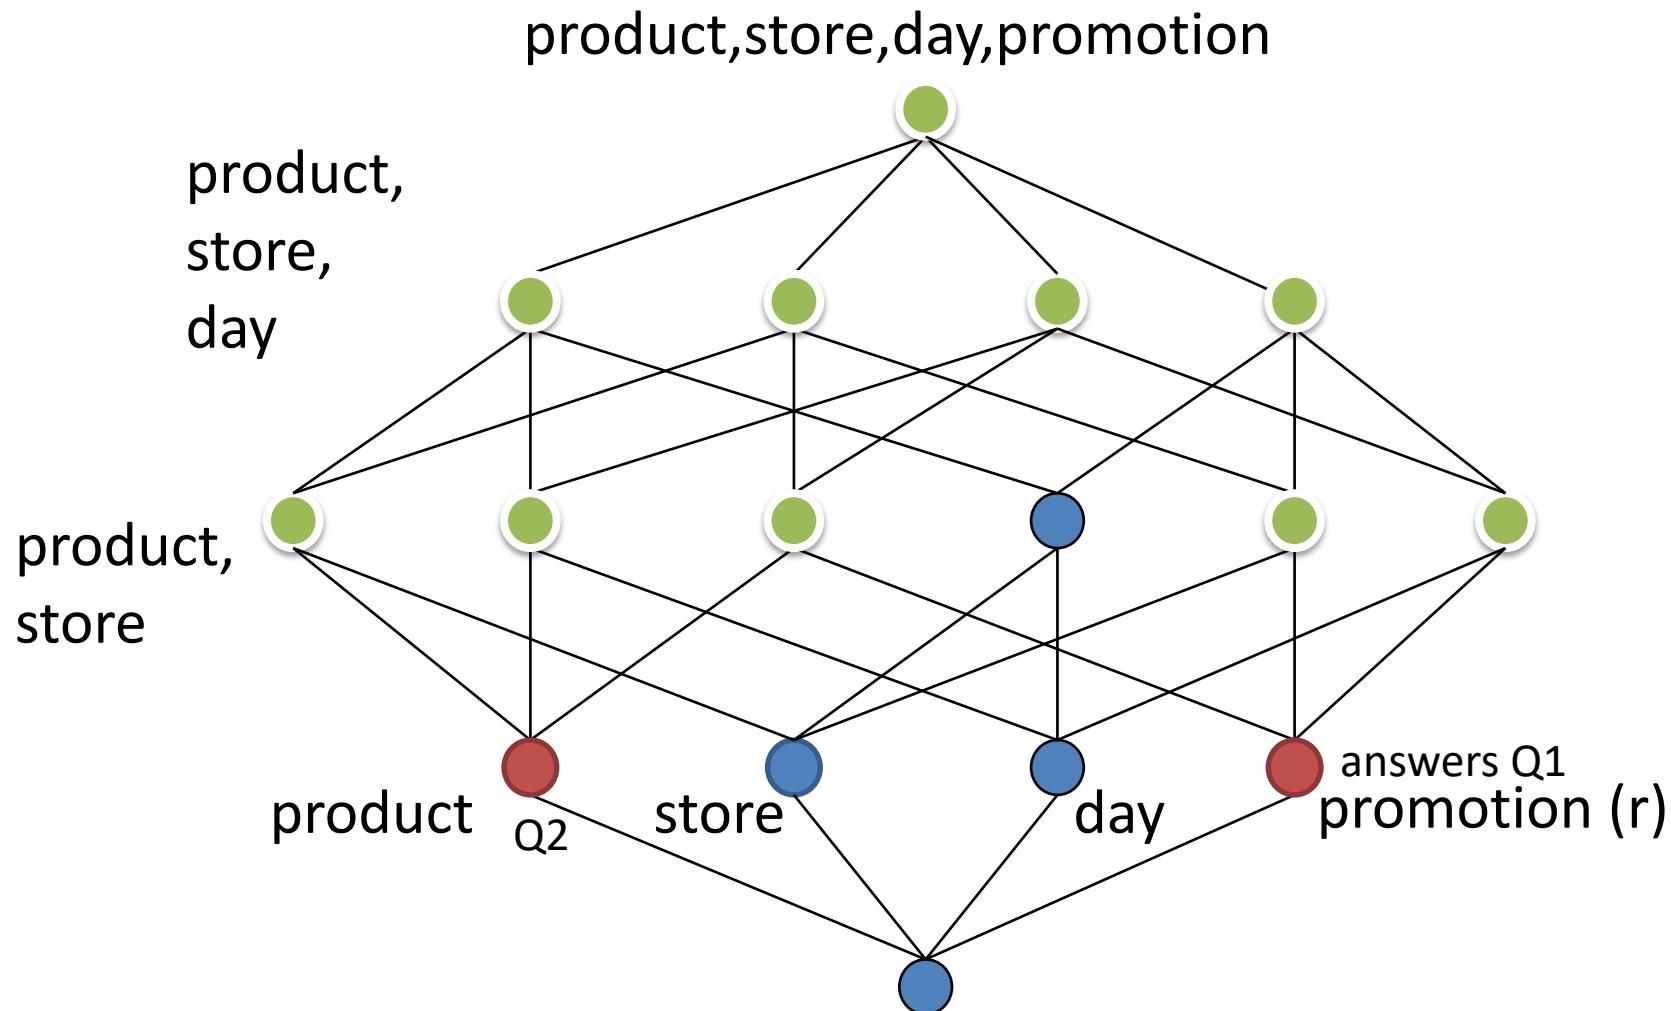


# Materialized View Selection Algorithm

1. Find lowest node for each  $Q$  in  $W$
2. Prune irrelevant sub-lattice
3. Evaluate cost function on any possible solution

Worst case :  $O(2^d)$  time     (  $d$  = #dimensions of the fact )

# Lattice has exponential size (in d)



# Materialized View Selection Algorithm

1. Find lowest node for each  $Q$  in  $W$
2. Prune irrelevant sub-lattice
3. Evaluate cost functions with an  
**approximated strategy**

PTIME approximation, but not guarantee of “optimality”.

# **INDEXING**

# Index Structures

- Traditional Access Methods
  - B+ trees
  - inverted lists
- Popular in Warehouses (for **dimensional** tables)
  - inverted lists
  - bit map indexes
  - join indexes

# The tuple acces problem

- Input
  - a set of attributes
  - and value for each attribtue (eg employée\_id = 17)
- Output
  - the tuples containg all input (attribute,value) pairs
    - sets of row\_id offsets
    - key IDs
- Index : a data structure that solves this problem in constant or logarithmic time

# B+ Trees

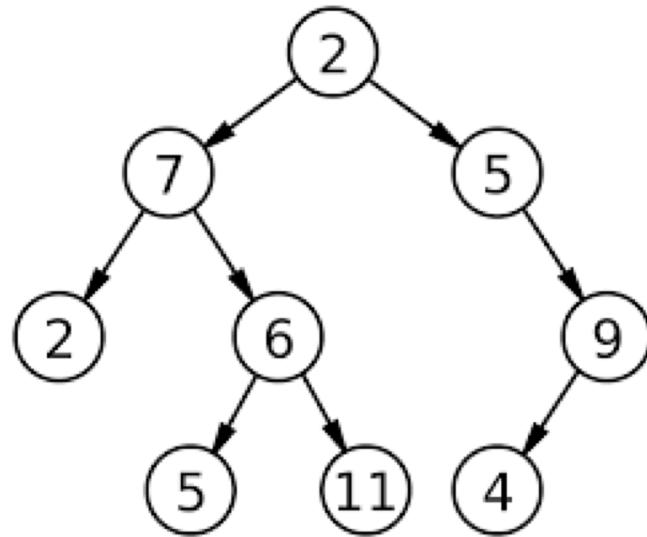
- Basic query optimization tasks must tackle queries like :

```
SELECT Employee.name  
      FROM Employee  
 WHERE Employee.id=17
```

id	name
1	Alice
2	Bob
3	Charle s
4	Bob
...	...
17	Dylan
...	...

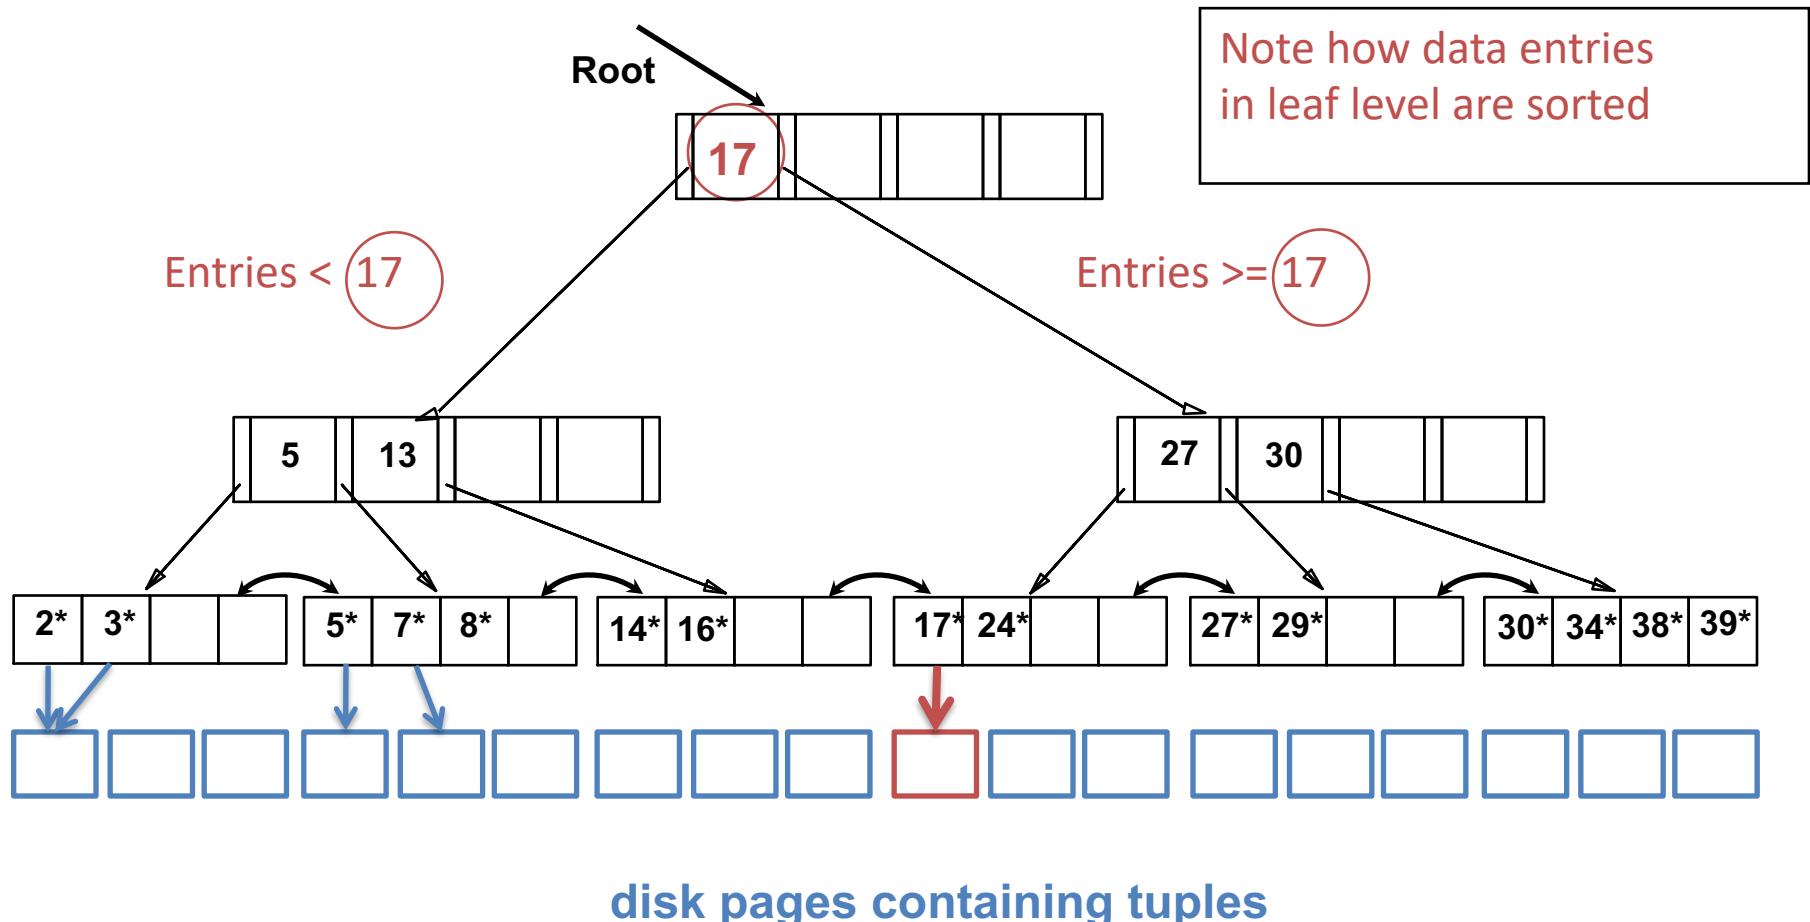
- access the tuple of employee 17 in constant time

# Binary trees are not enough



- reduce deepnesses → cut index access time cost

# B+ Tree on Employee.id



# Inverted Lists

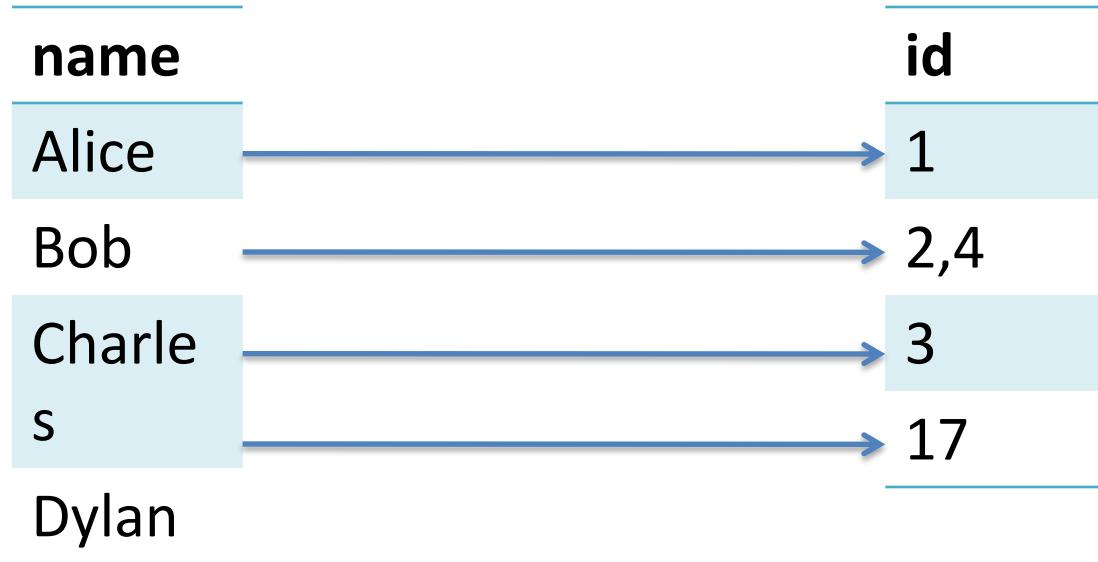
- Basic query optimization tasks on datawarehouse dimensions must tackle queries like :

```
SELECT Employee.id  
      FROM Employee  
     WHERE name="Bob"
```

id	name
1	Alice
2	Bob
3	Charle s
4	Bob
...	
17	Dylan
...	

- access all tuples in constant time

# Inverted Lists



# Combining Inverted Lists

- Query:
  - Get people with age = 20 and name = “Bob”
- List for age = 20
  - r2, r5, r18
- List for name = “Bob”
  - r2, r4
- Answer is intersection: r2

# Bit Map Index

- Alternative representation of RecordID-list

# Bitmap Index Example

<b>custid</b>	<b>name</b>	<b>gender</b>	<b>rating</b>
112	Joe	M	3
115	Ram	M	5
119	Sue	F	5
116	Woo	M	4

Query : Find all customers with gender='M'

# Bitmap Index Example

	<b>custid</b>	<b>name</b>	<b>gender</b>	<b>rating</b>
true	112	Joe	M	3
true	115	Ram	M	5
false	119	Sue	F	5
true	116	Woo	M	4

Query : Find all customers with gender='M'

# Bitmap Index Example

gender=M
1
1
0
1

true  
true  
false  
true

custid	name	gender	rating
112	Joe	M	3
115	Ram	M	5
119	Sue	F	5
116	Woo	M	4

Query : Find all customers with gender='M'

# Bitmap Index Example

gender=F
0
0
1
0

false  
false  
true  
false

custid	name	gender	rating
112	Joe	M	3
115	Ram	M	5
119	Sue	F	5
116	Woo	M	4

Query : Find all customers with gender='F'

# Bitmap Index Example

rating=5				
0	←	false	custid	name
1	←	true	112	Joe
1	←	true	115	Ram
0	←	false	119	Sue
			116	Woo

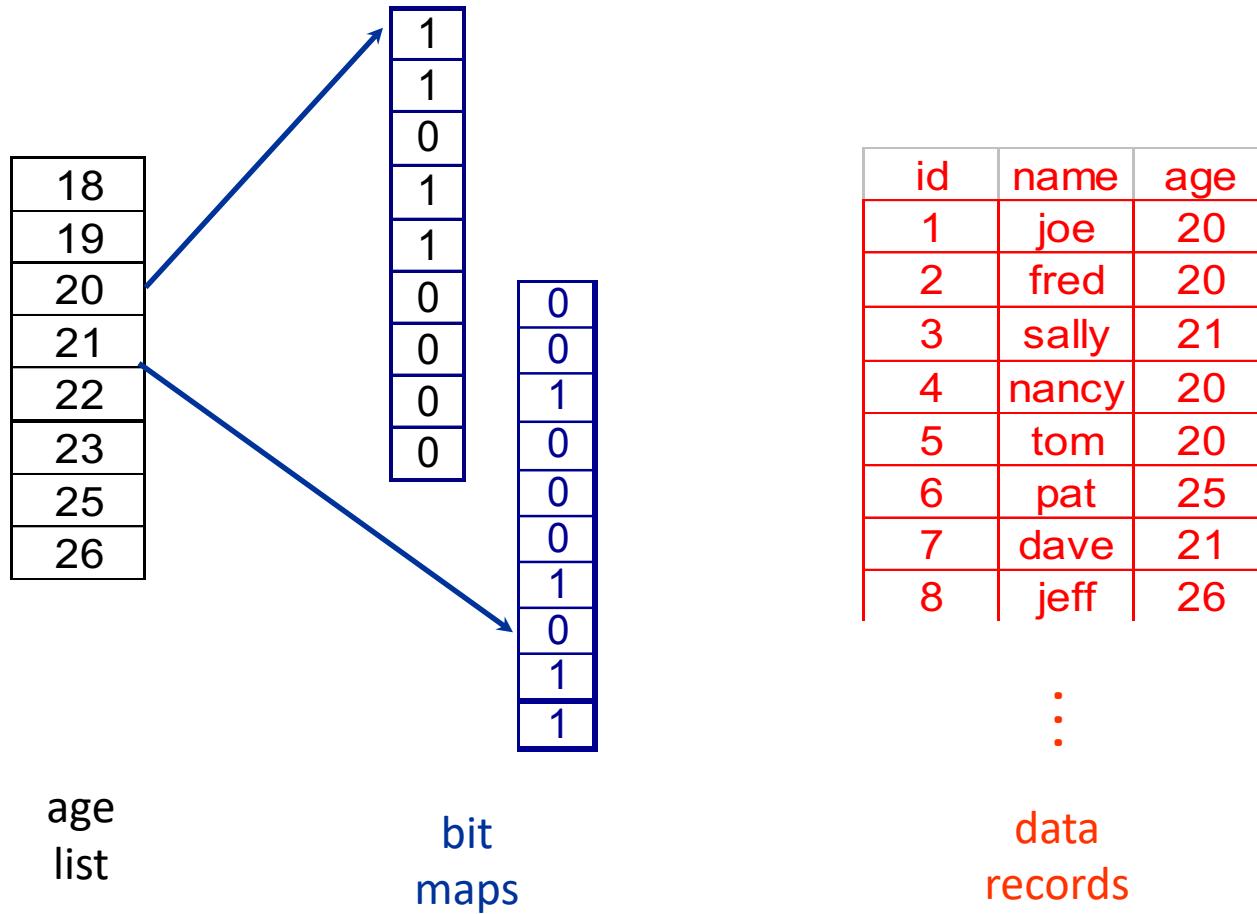
Query : Find all customers with rating = 5

# Bitmap Index Example

custid	name	gender	rating
112	Joe	M	3
115	Ram	M	5
119	Sue	F	5
116	Woo	M	4

rating=1	rating=2	rating=3	rating=4	rating=5
0	0	1	0	0
0	0	0	0	1
0	0	0	0	1
0	0	0	1	0

# Bit Maps



# The magic of Booleans

- Query:
  - Get people with age = 20 and name = “Bob”
- List for age = 20: **1101100000**
- List for name = “Bob”: **0100000001**
- Answer is logical-& **0100000000**

# The magic of Booleans

- Query:
  - Get people with age = 20 OR age = 21
- List for age = 20: **1101100000**
- List for age = 21: **0010000001**
- Answer is logical-OR **111100001**

# Bitmap indexes

- Advantageous for **low-cardinality domains**
  - Significant reduction in space and I/O  
Compression (e.g., run-length encoding) exploited
- Comparison, join and aggregation operations  
are reduced to *bit arithmetic*

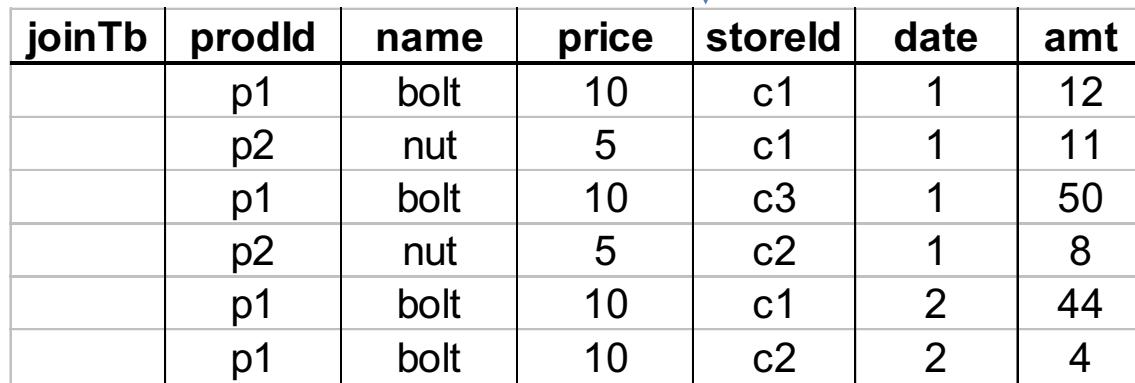
## Joins

## Fact\_table\_sales

sale	prodId	storeId	date	amt
	p1	c1	1	12
	p2	c1	1	11
	p1	c3	1	50
	p2	c2	1	8
	p1	c1	2	44
	p1	c2	2	4

## Dim\_table\_product

product	id	name	price
	p1	bolt	10
	p2	nut	5



# Join Indexes

Dim\_table\_product

join index

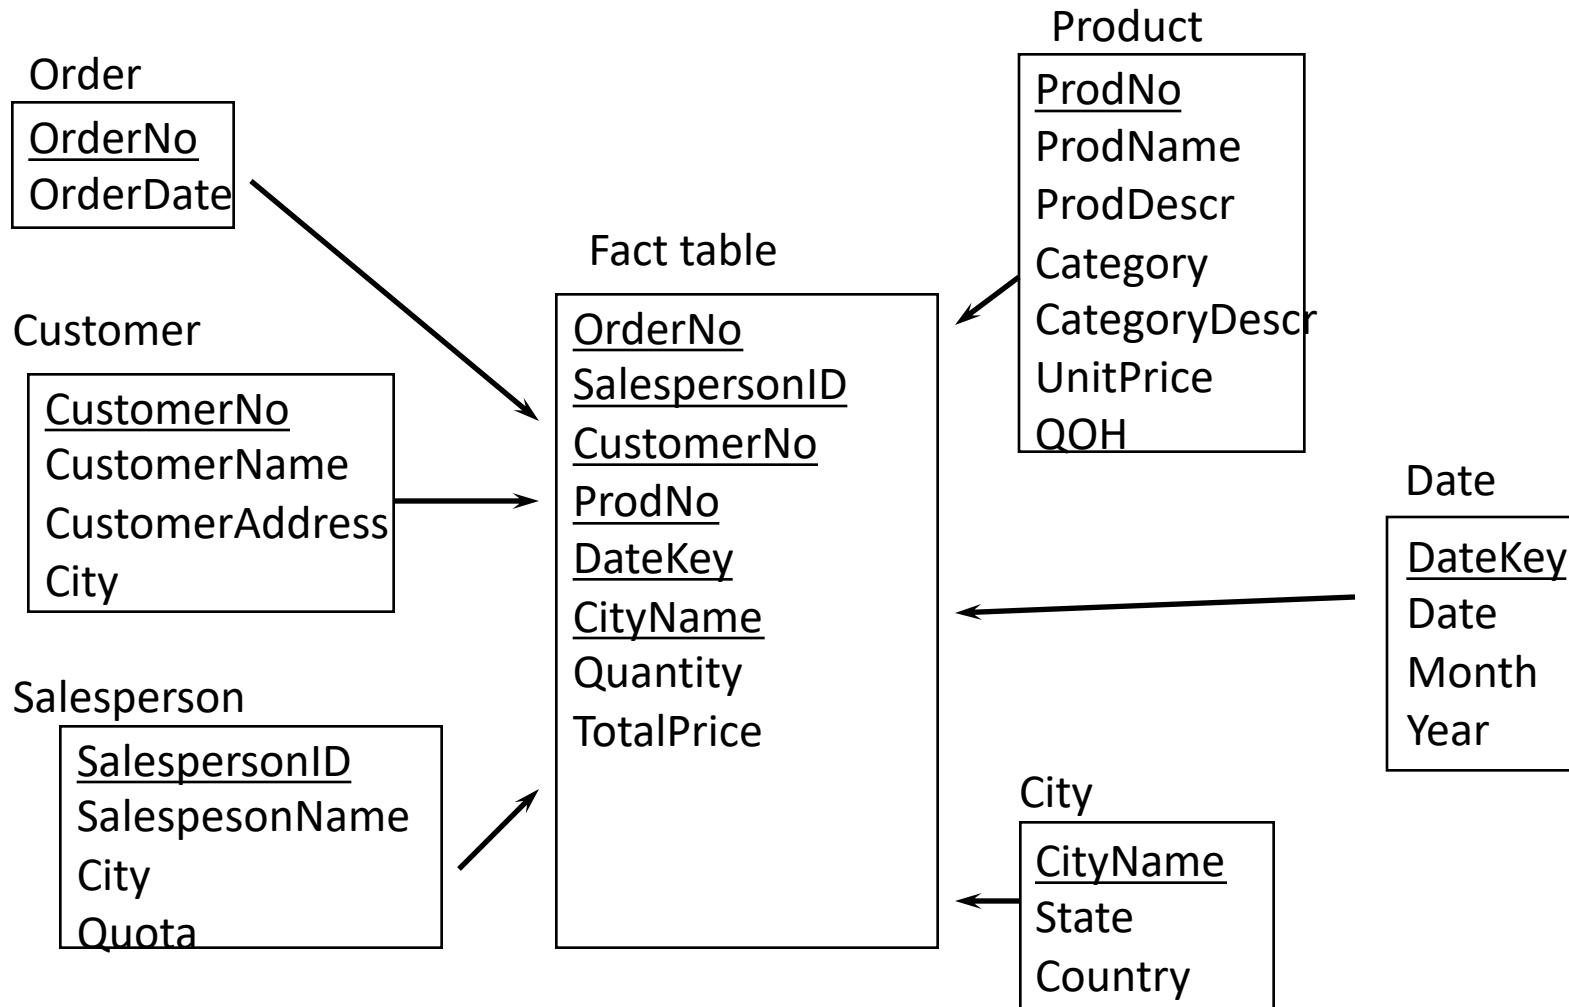
product	id	name	price	jIndex
	p1	bolt	10	r1,r3,r5,r6
	p2	nut	5	r2,r4

sale	rId	prodId	storeId	date	amt
	r1	p1	c1	1	12
	r2	p2	c1	1	11
	r3	p1	c3	1	50
	r4	p2	c2	1	8
	r5	p1	c1	2	44
	r6	p1	c2	2	4

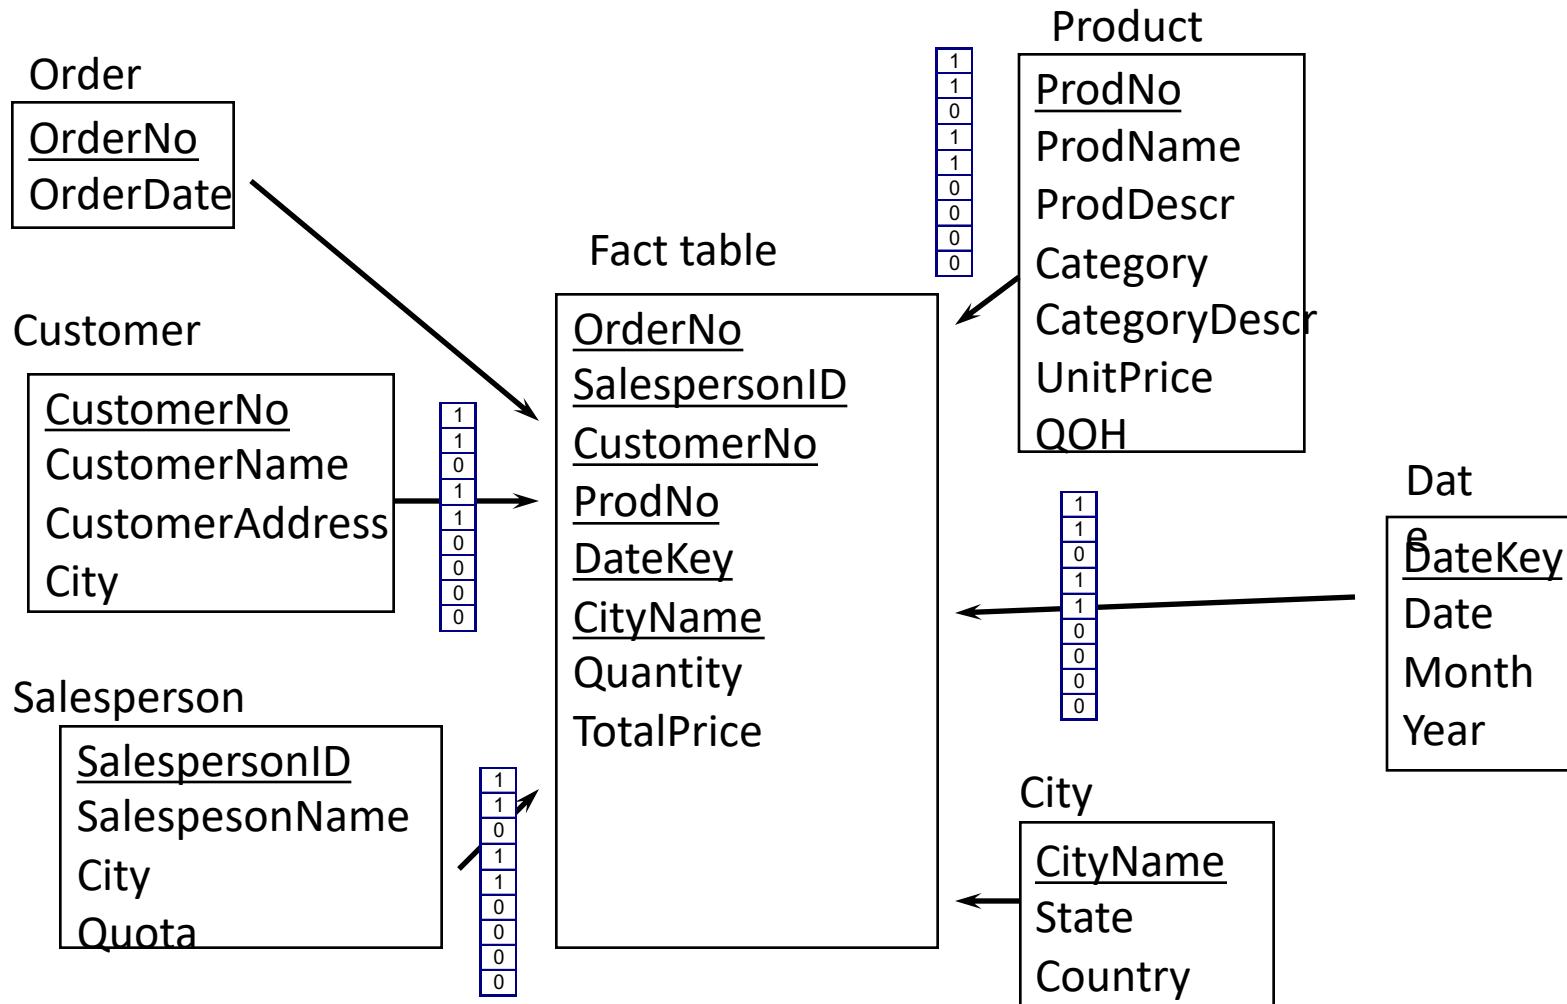
# Join Index

- Relationships between attribute value of a dimension and the matching rows in fact table
- Good news : it can be bitmapped !

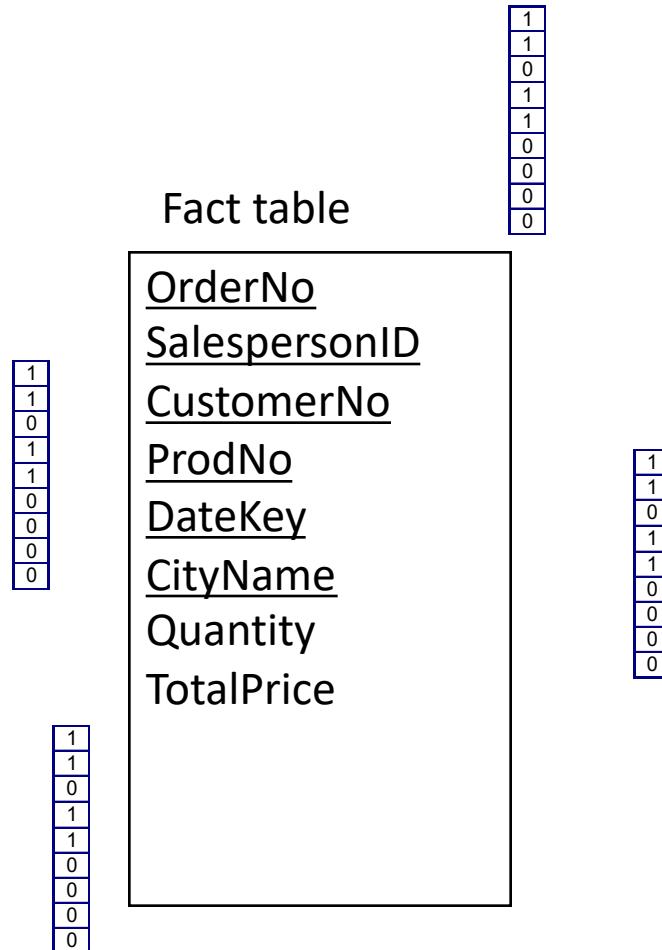
# Join Index over Star Schema



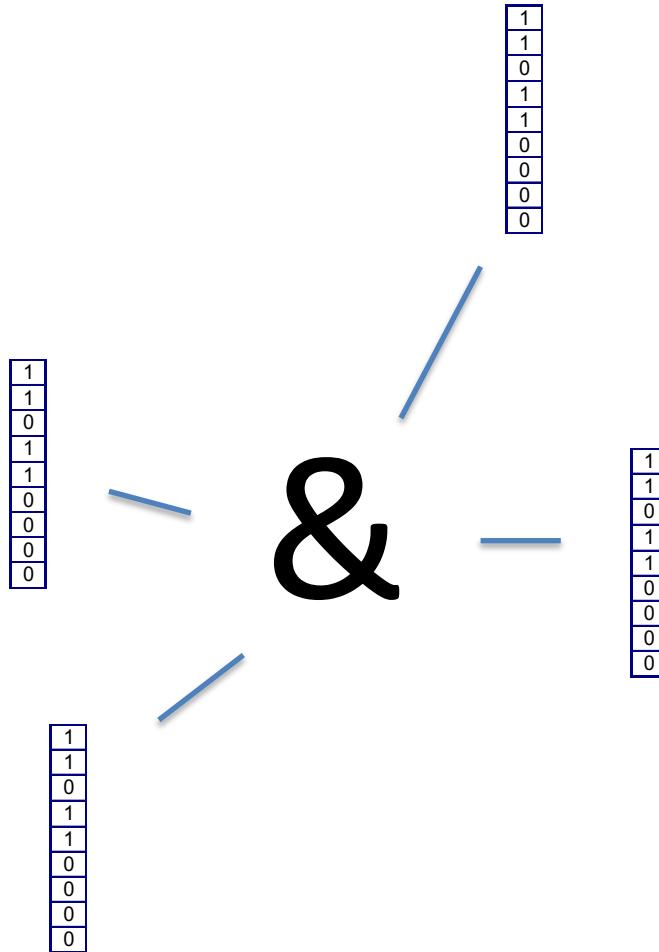
# Join Index over Star Schema



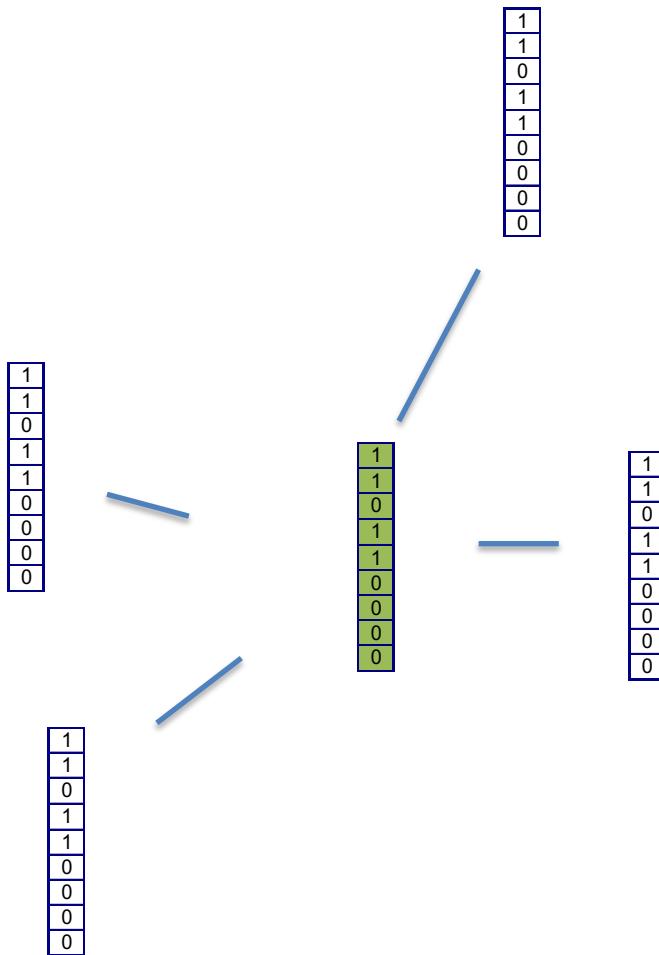
# Join Index over Star Schema



# Join Index over Star Schema



# Join Index over Star Schema

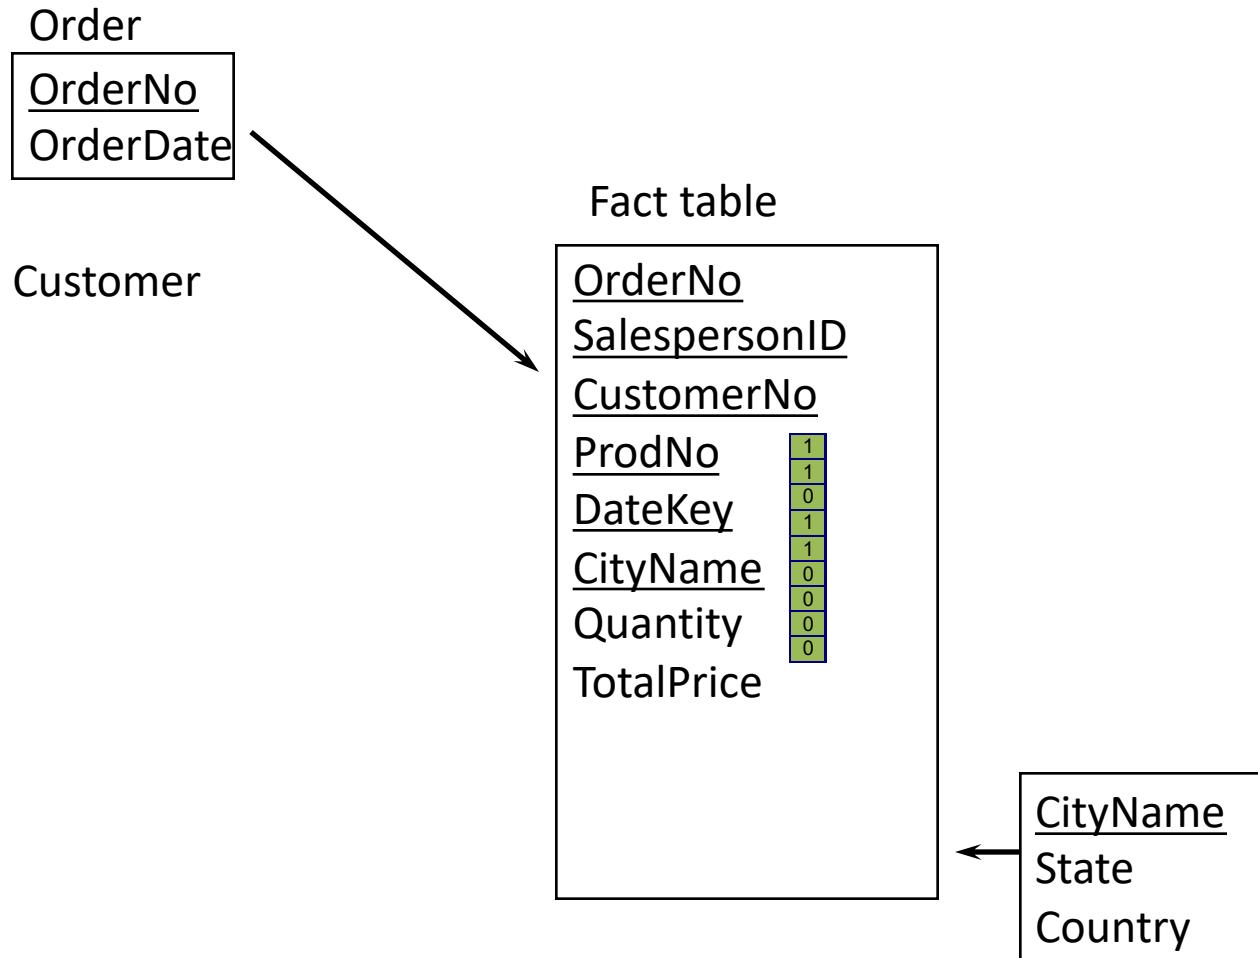


# Join Index over Star Schema

Fact table

<u>OrderNo</u>	
<u>SalespersonID</u>	
<u>CustomerNo</u>	
<u>ProdNo</u>	1 1 0 1 1
<u>DateKey</u>	
<u>CityName</u>	0 0 0 0
<u>Quantity</u>	
<u>TotalPrice</u>	

# Join Index over Star Schema



# Join-Algorithm Using Bitmapped Join-Indexes

1. Take intersection of join-indexes for all available dimensions so as to filter as many facts as possible
2. Join the remaining dimension on the resulting facts
3. Aggregate on the fact table

# MULTIDIMENSIONAL DATABASES

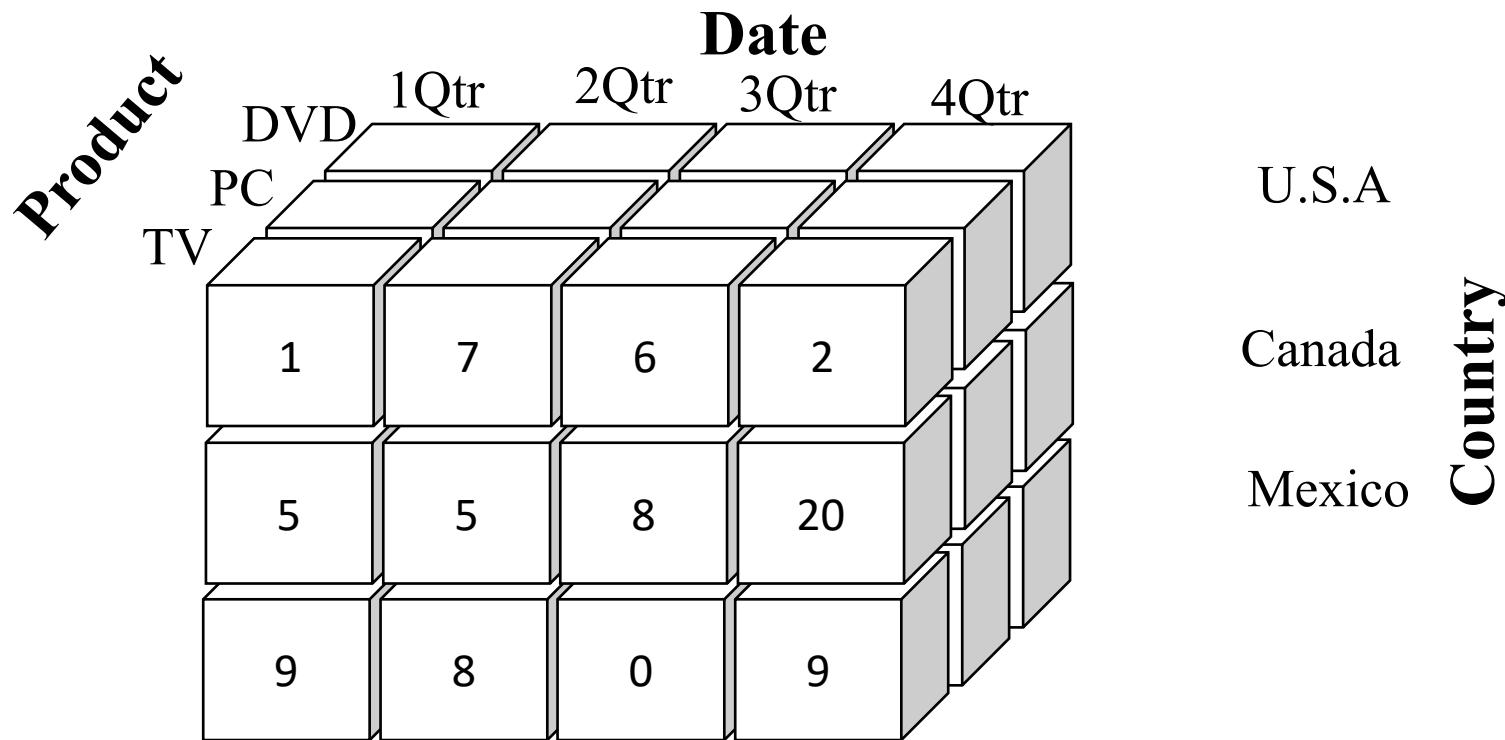
# From Tables to (hyper-)Cubes

- Datamarts are **multidimensional models** which views data in the form of a data (hyper-)cube
- We have seen relational datawarehouses, just one kind of “implementation” of this model
- Other implementations, may differ in the technological solution
  - for example, how to store sparse arrays (lots of zeros) with a matrix-based approach

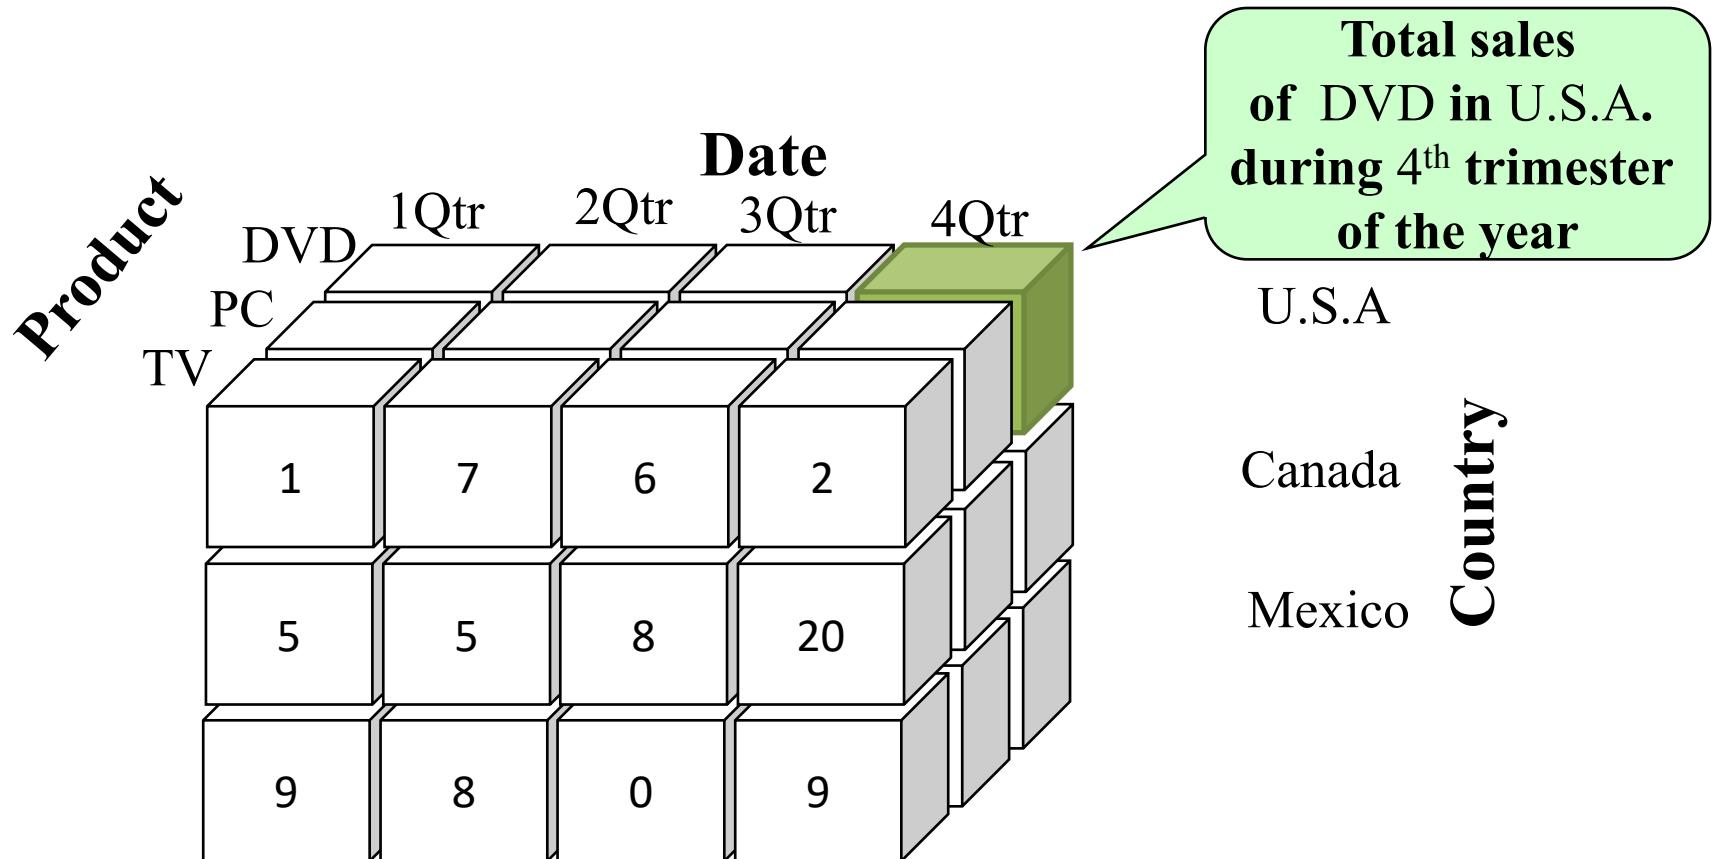
# From Tables to (hyper-)Cubes

- A data cube, such as `sales`, allows data to be modeled and viewed in multiple dimensions
  - Dimension tables, such as  
`item (item_name, brand, type)`, or  
`time(day, week, month, quarter, year)`
  - Fact table contains measures (such as  
`dollars_sold`) and keys to each of the related dimension tables

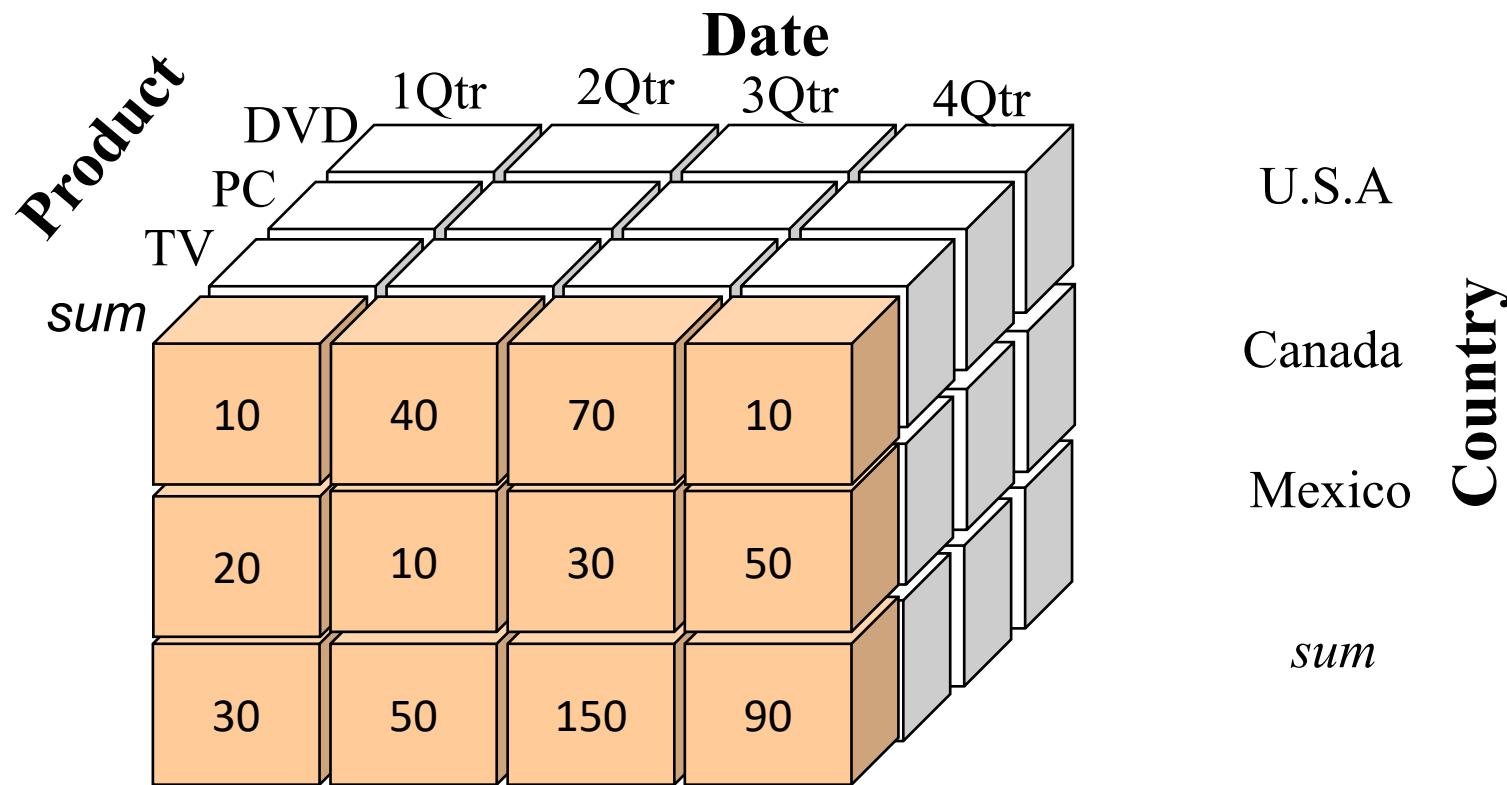
# A Sample Data Cube



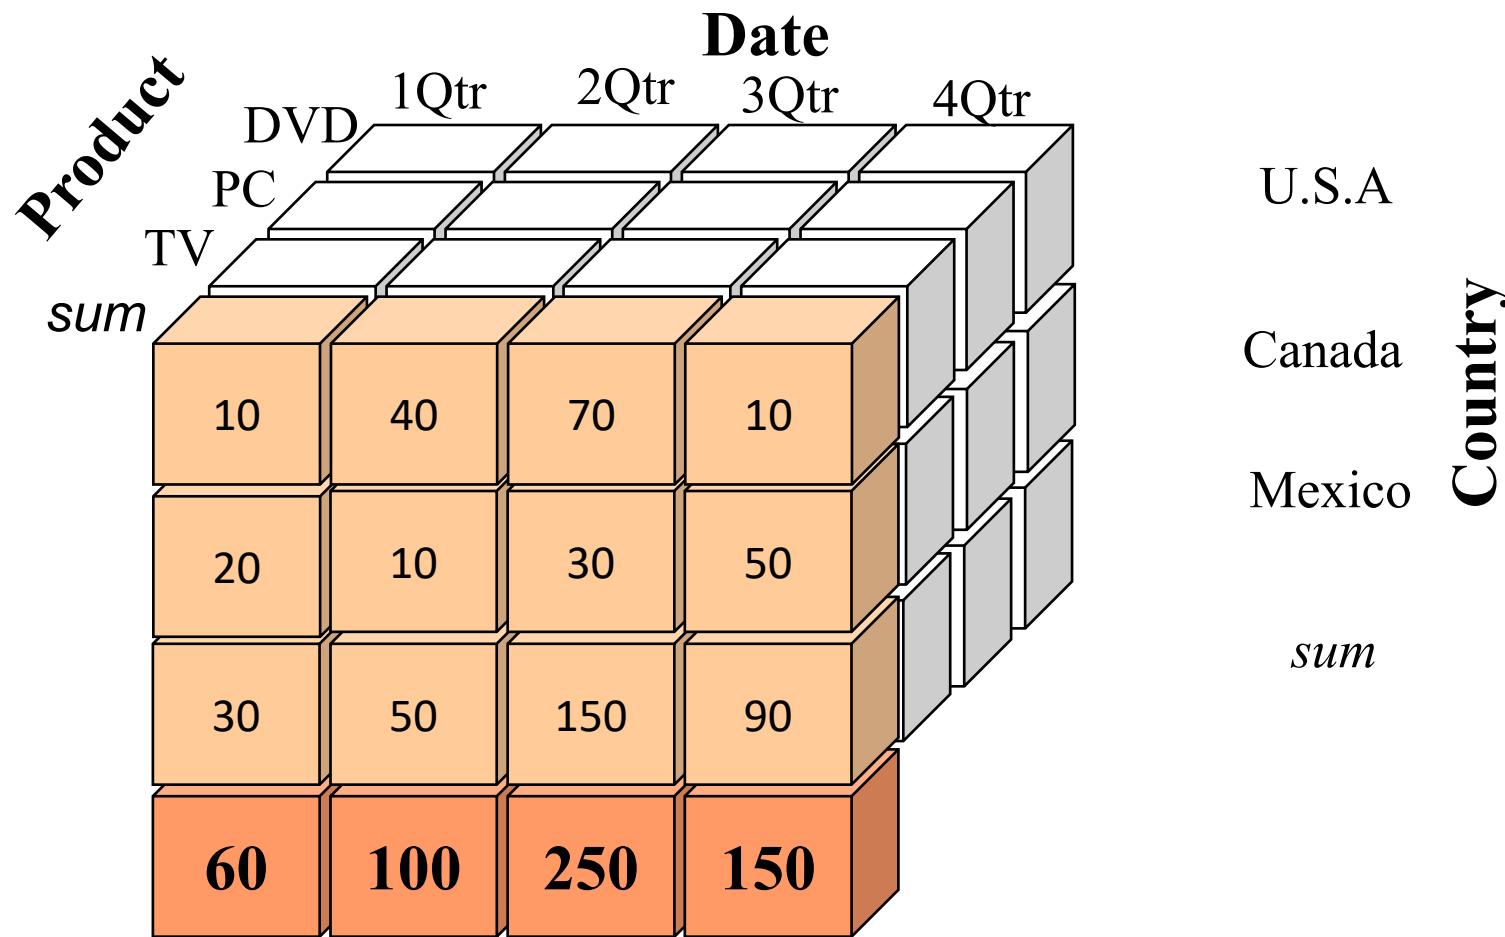
# A Sample Data Cube



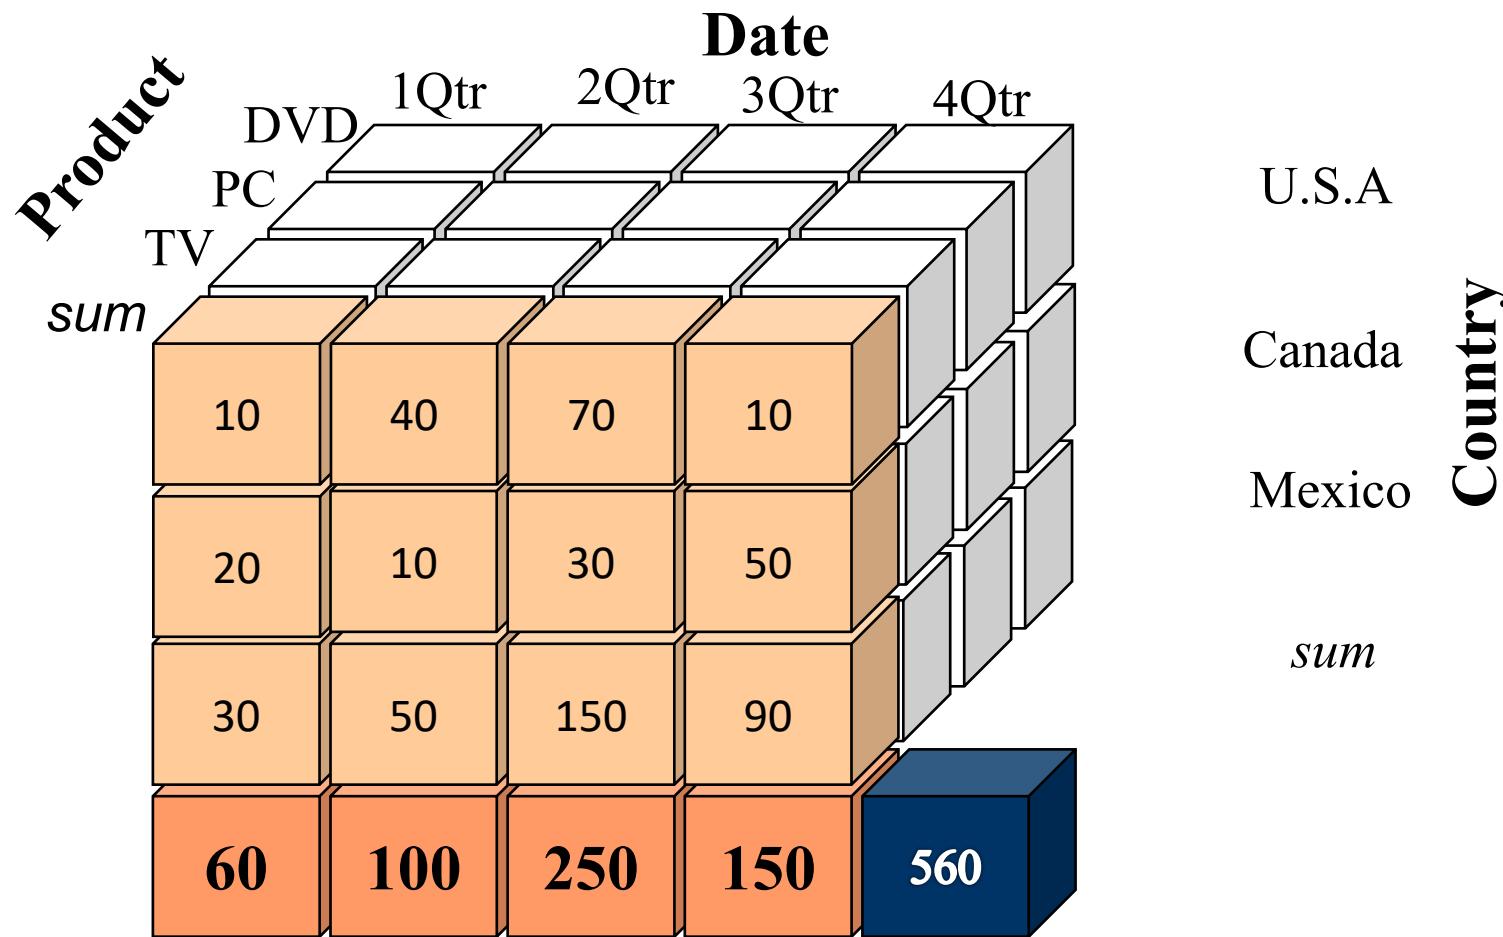
# Group By Date, Country



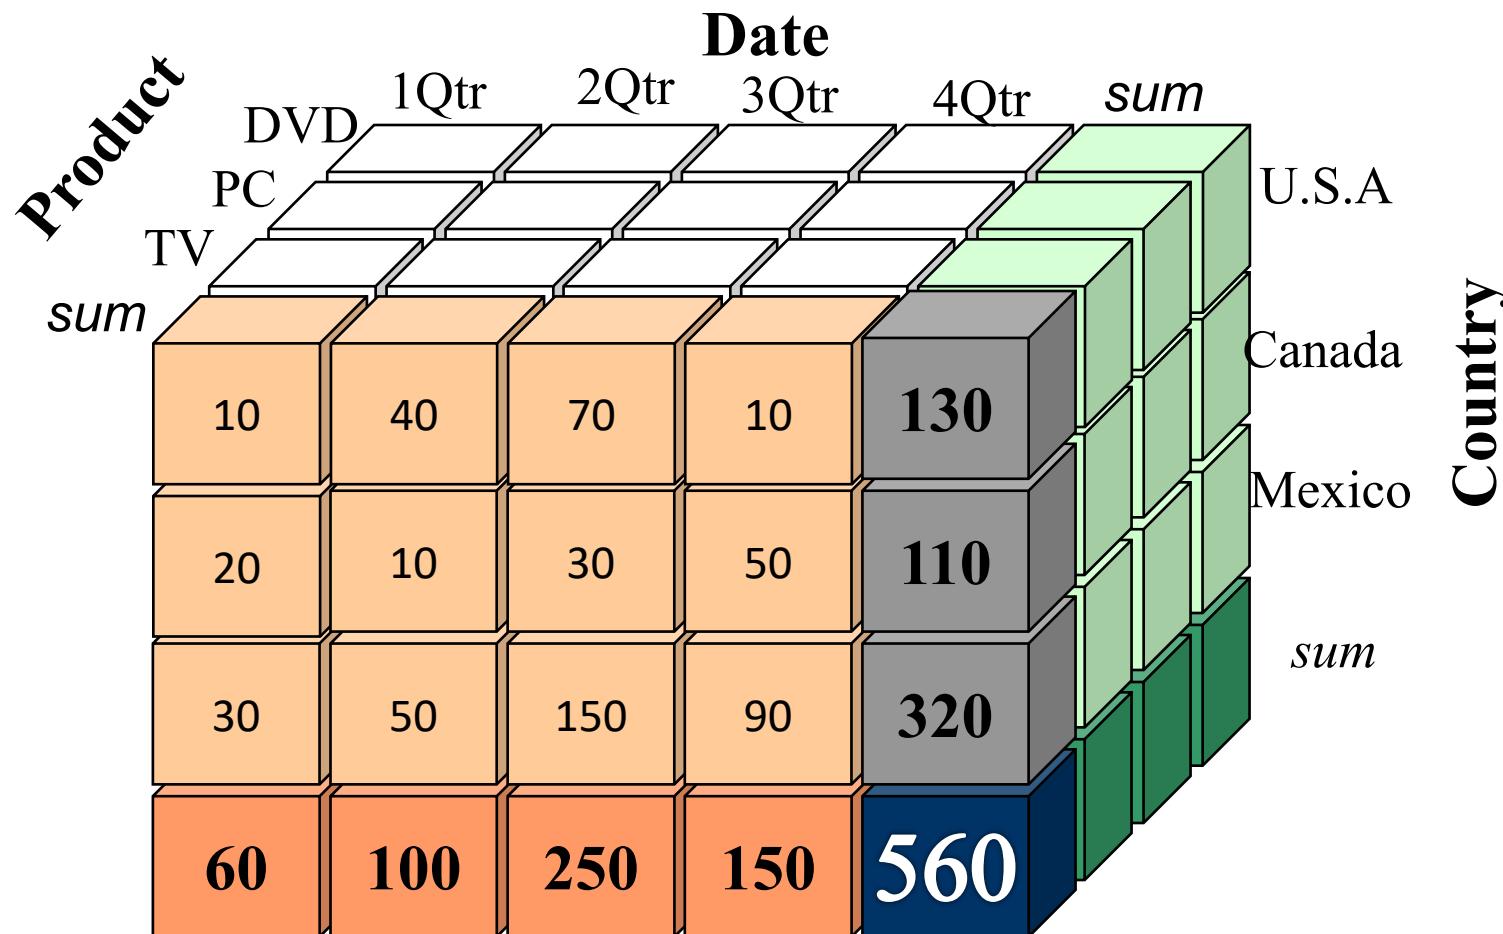
# Group By Date



# Group By ROLLUP (Date, Country)



# Group By CUBE (Period, Date, Country)



# Typical OLAP Operations

- **Roll up** (drill-up): summarize data
  - *by climbing up hierarchy or by dimension reduction*
- **Drill down** (roll down): reverse of roll-up
  - *from higher level summary to lower level summary or detailed data, or introducing new dimensions*
- **Slice and dice:**
  - *project and select*
- **Pivot (rotate):**
  - *reorient the cube, visualization, 3D to series of 2D planes.*
- Other operations
  - **drill across:** *involving (across) more than one fact table*
  - **drill through:** *through the bottom level of the cube to its back-end relational tables*

GROUP-BY  
ROLL-UP  
CUBE

# Multidimensional DB History

- '90 large datawarehouses appeared
- '98 MS OLAP Server, first multidim. database for mass market
- Today: multidimensional analytics part of any commercial DB
- 2000 : Column-stores as the novel data warehouse systems
- 2010+ : Big Data : exploratory analytics